


Introduction to programming paradigms

August 15, 2013

Universidad del Valle
AVISPA Research Group

Peter Van Roy
Université catholique de Louvain
Louvain-la-Neuve, Belgium

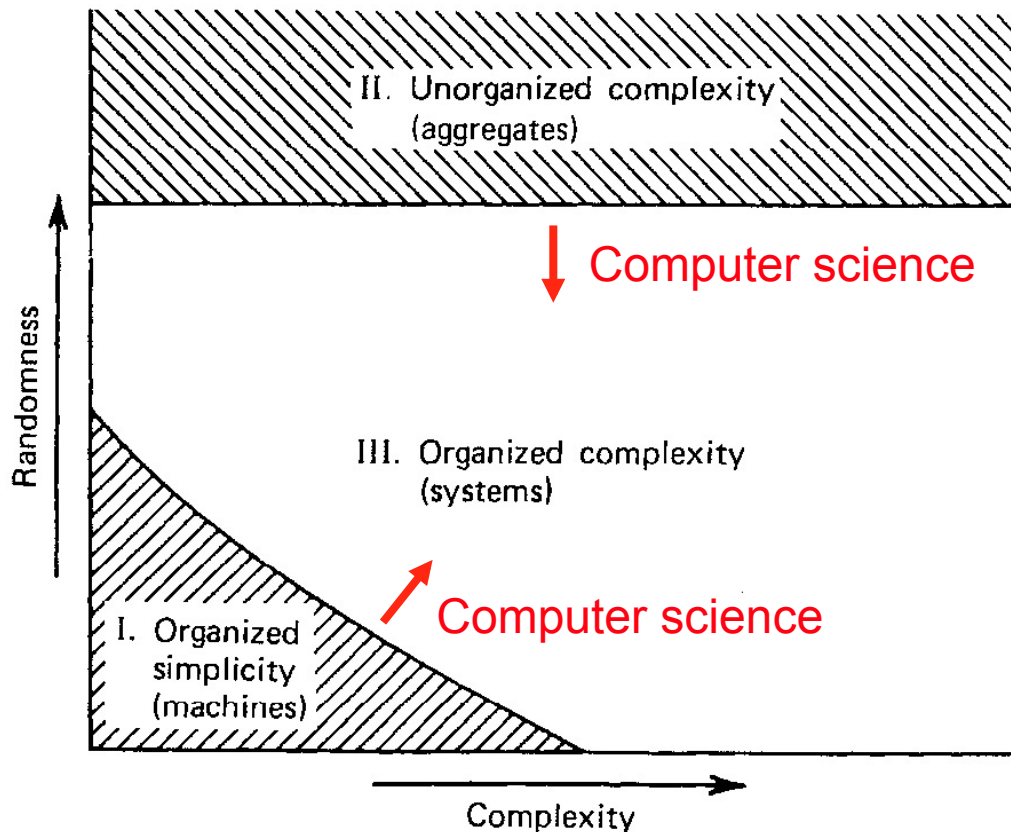




Programming paradigms

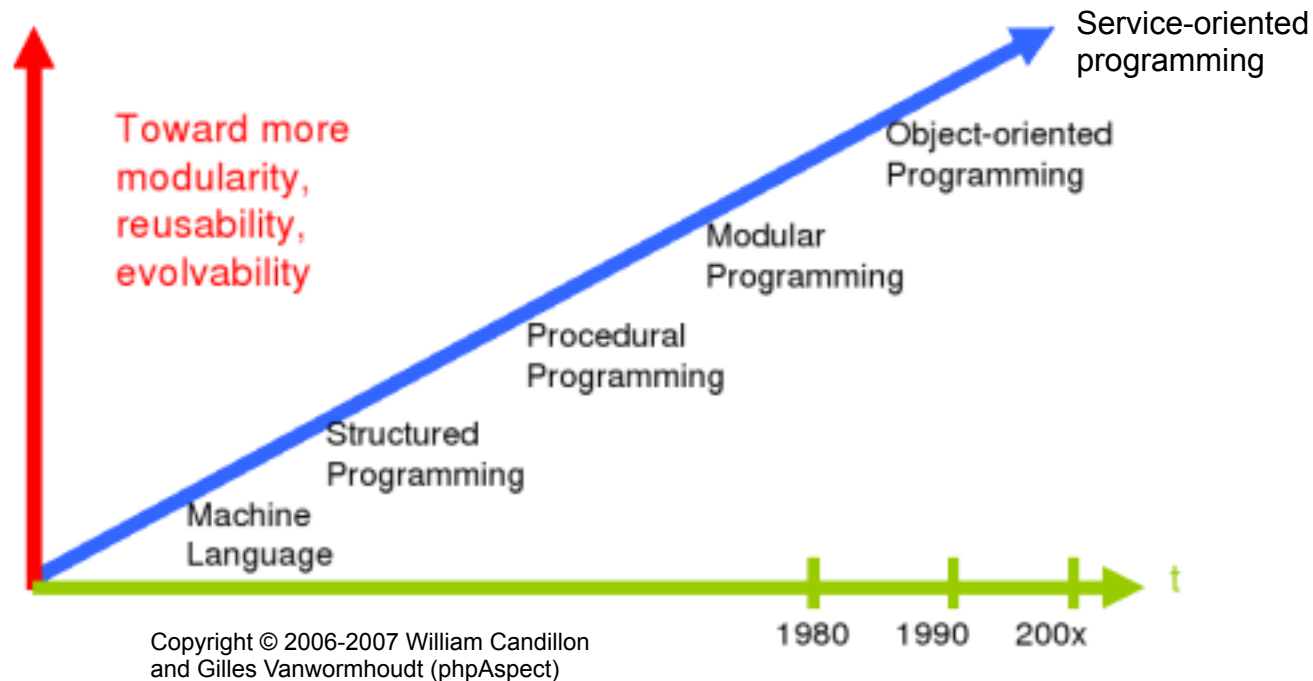
- A **paradigm** is *an approach to programming a computer based on a coherent set of principles or a mathematical theory*
 - Different theories of computing result in different paradigms (λ calculus, π calculus, first-order logic, Hoare logic, ...)
 - No existing theory covers all programming concepts!
 - Programming is truly **a new discipline** that is not covered by traditional mathematical theories
- We will introduce the world of programming paradigms
 - Why do we need many paradigms?
 - Because solving a problem is **much easier** when done in the right paradigm!

Programming and complex systems




- **Systems** are composed of parts that interact in a well-defined way to provide a new behavior
 - This diagram comes from [Weinberg 1977] *An Introduction to General Systems Thinking*
- Computer science is the most advanced discipline for building complex systems
- Programming paradigms are the vanguard of building complex systems
- Each newly discovered programming concept advances the science of building complex systems

Traditional view of programming progress (1)



- This is a typical diagram illustrating how programming has progressed
- This diagram only says a little bit; it leaves out many important ideas

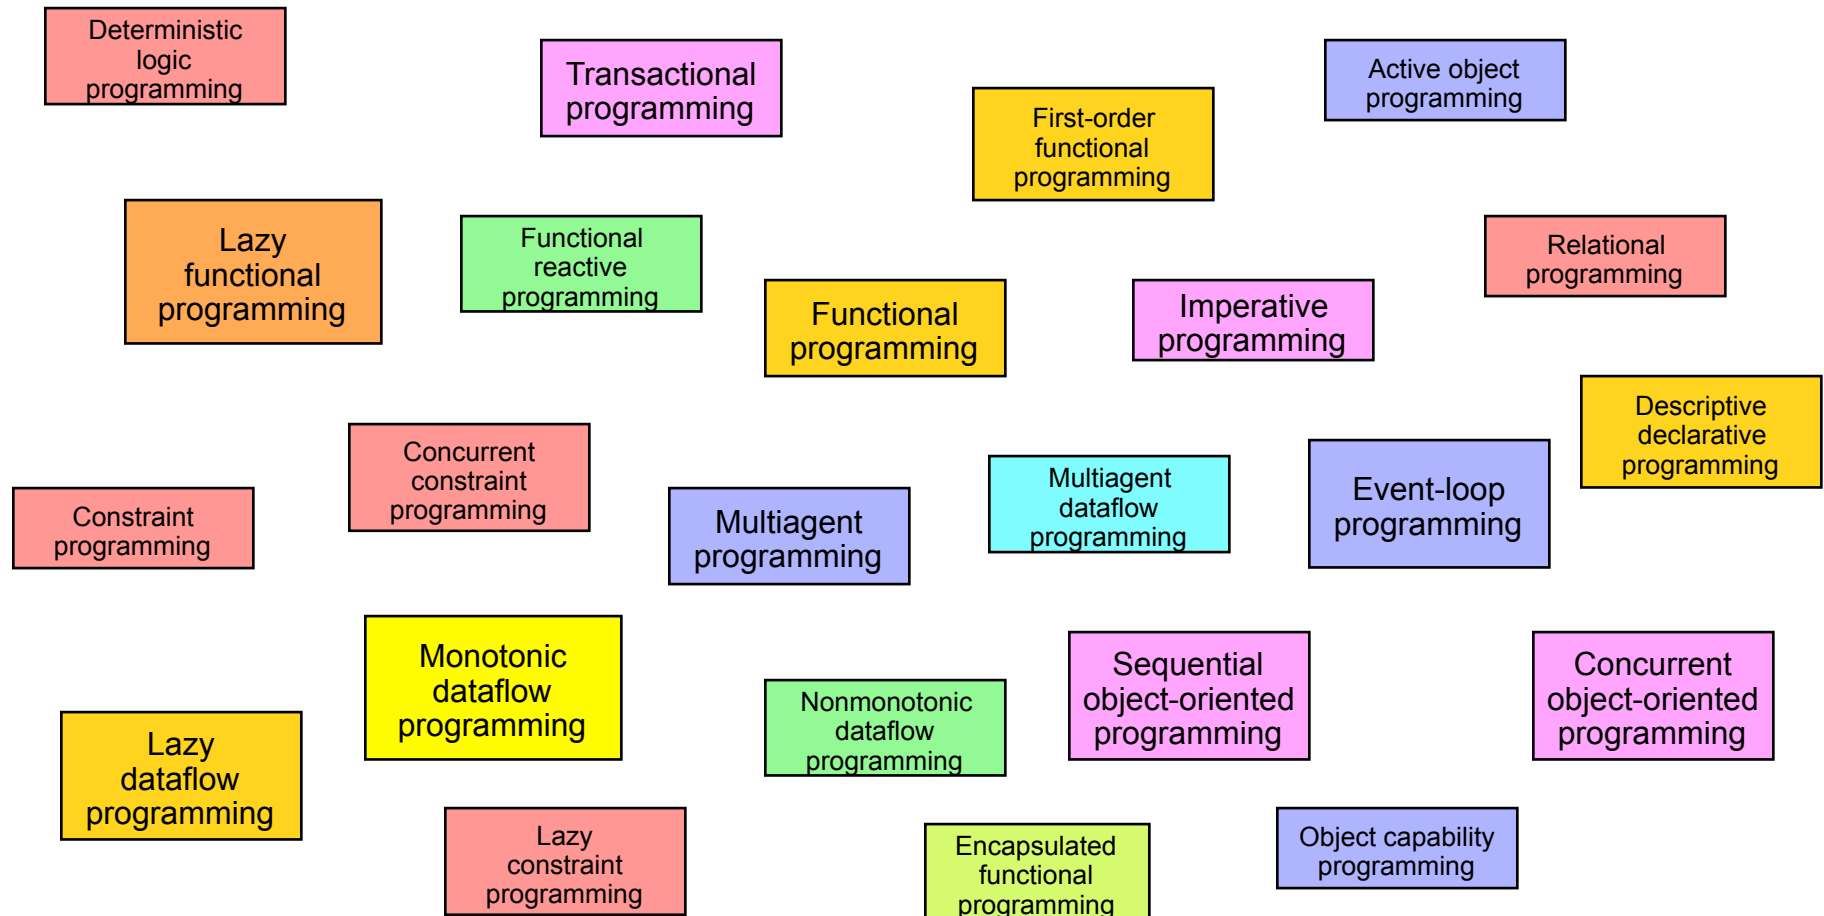

Traditional view of programming progress (2)




Copyright © 2004 Christophe de Dinechin (SourceForge.net)

- This diagram is more realistic
- Each new paradigm simplifies programming: it raises the level of complexity that is easy to program
- Again, this diagram does not take into account many important programming concepts
 - At least, distributed (networked) programming is mentioned, which is part of the paradigm of concurrent programming!

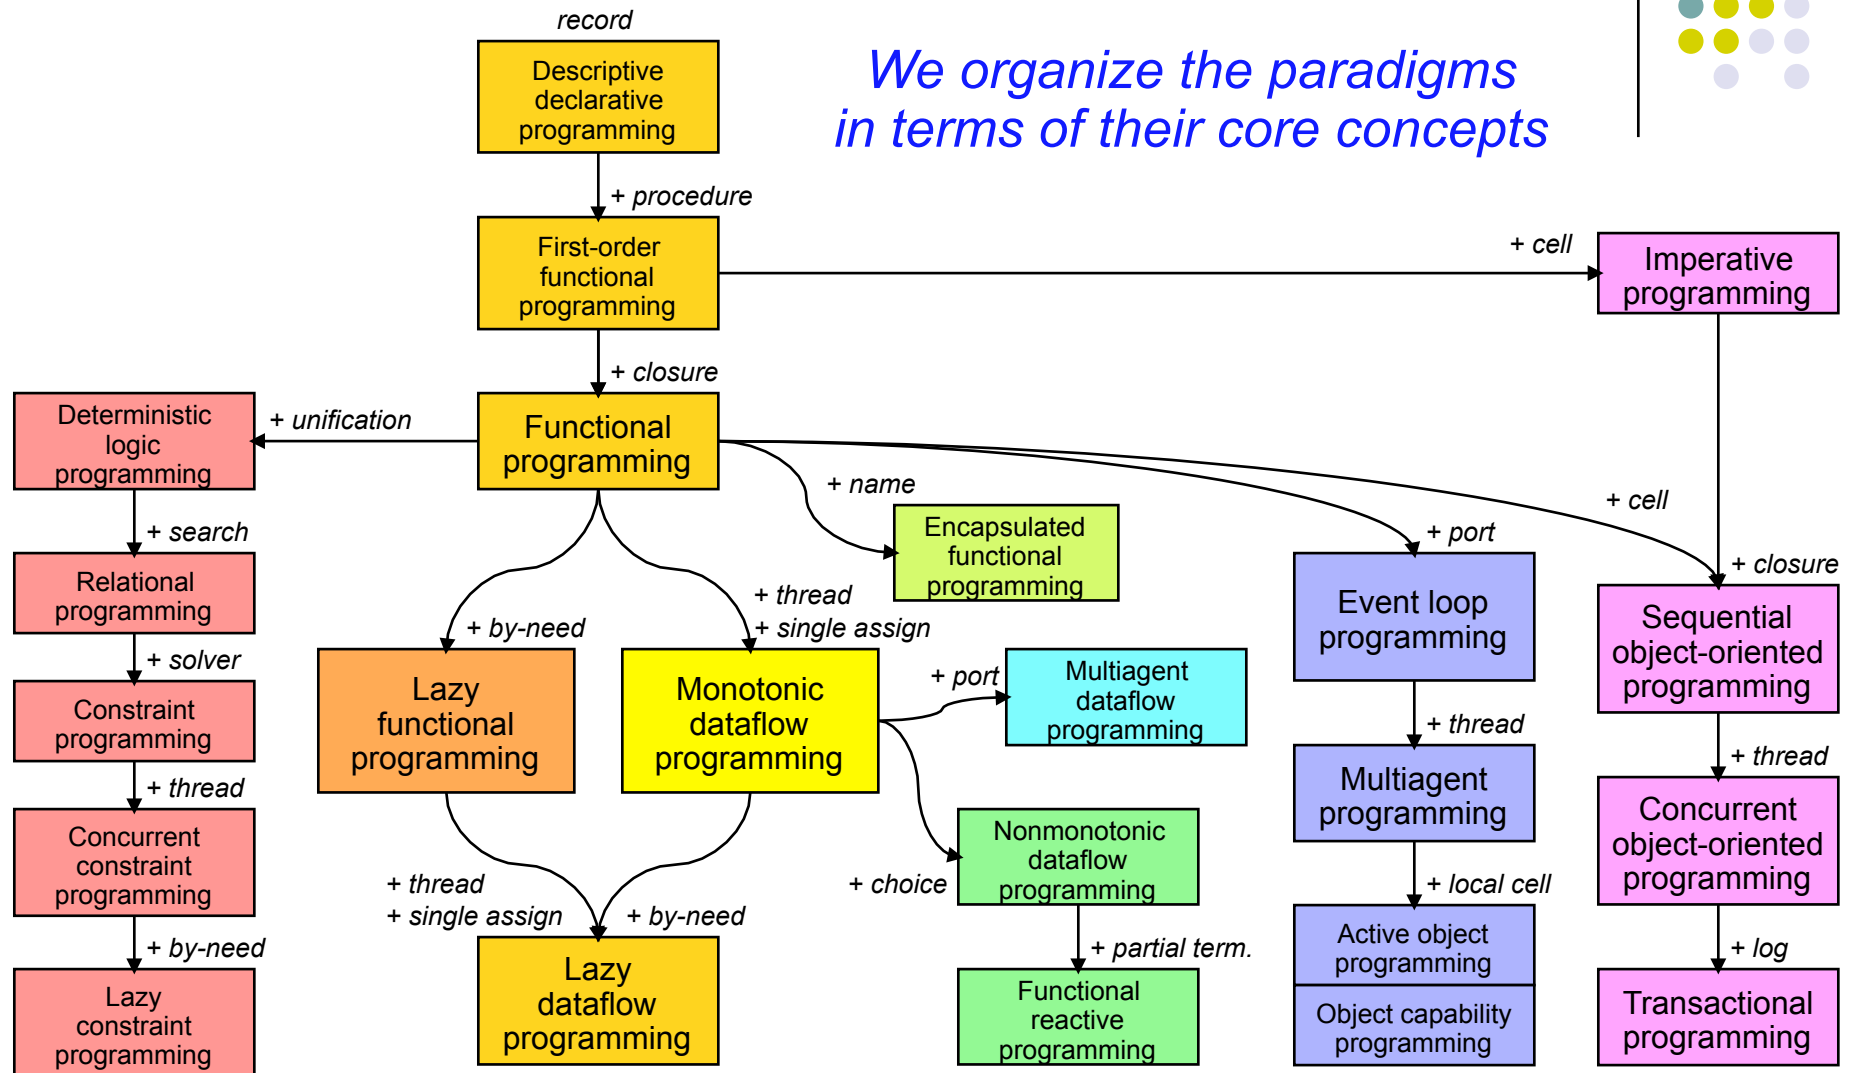
The paradigm jungle




Taxonomy of paradigms




We organize the paradigms in terms of their core concepts



Concurrency and nondeterminism





Concurrency


- The real world is **concurrent**
 - It consists of **activities that evolve independently**
- The computing world is concurrent too, at three levels
 - Distributed system: computers linked by a network
 - A concurrent activity is called a **computer**
 - Operating system of one computer
 - A concurrent activity is called a **process**
 - Each process has an independent memory space (competitive)
 - Activities inside one process
 - A concurrent activity is called a **thread**
 - All threads share the same memory space (cooperative)



Concurrency is really hard?

- At least, that's what people believe!
- It's true in C++ and Java: synchronized objects (monitors) are really hard to program with
- But the problem is not concurrency, it's monitors!
 - Monitors are the **wrong paradigm for concurrency**
- So what's the right paradigm?
 - That's what we'll see now!

The right paradigm: deterministic dataflow



- Producer/consumer pipeline

```
fun {Prod N Max}  
  if N < Max then  
    N | {Prod N+1 Max}  
  else nil end  
end
```




```
proc {Cons Xs}  
  case Xs of X|Xr then  
    {Display X}  
    {Cons Xr}  
  [] nil then skip end  
end
```


```
local Xs in  
  thread Xs = {Prod 0 1000} end  
  thread {Cons Xs} end  
end
```

- The threads in Prod and Cons share the dataflow variable **Xs**
- Dataflow behavior of the **case** (wait until data arrives) results in *stream communication*
- No other synchronization is needed!

Is deterministic dataflow enough?



- Unfortunately not, sometimes it's not expressive enough
- Simple example: two clients talking to one server




C1 and C2 are independent: requests arrive at S in any order. If two requests arrive at the same time, S has to choose one of them.

This cannot be programmed with deterministic dataflow! S cannot receive messages in a fixed order because S does not know which client will send a message next. S must handle the messages as they arrive.



Nondeterministic choice

- Nondeterministic choice appears when two or more entities (clients) interact with the same entity (server)
 - Each client interacts with the server independently of the other clients
- The server must decide which message to handle first
 - The server (the runtime system, not the application) makes a **choice**
 - This choice is called **nondeterminism**
- Nondeterministic choice is **a new programming concept**
- Adding it gives a new paradigm, **multiagent dataflow programming** (which is similar to the actor model or message passing)



A few small examples...


- Deterministic dataflow

- Digital logic simulation (clocked execution)
- Josephus problem (ring protocol)
- Hamming problem (lazy protocol with cycle)
- Bounded buffer (combining lazy and eager)

+ nondet.
choice

- Multiagent dataflow

- Multiagent systems (agents talking to each other)
- Distributed protocols (RMI and distributed algorithms)
- Lift control system (real-world example multiagent system)



A bigger example...


- Ericsson AXD 301 ATM Switch: >1 million lines of Erlang

- **Erlang**: Concurrent and independent by default, asynchronous messages, multi-agent programs

- **Java**: Sequential and monolithic by default, synchronous RMI, shared-data programs

- Object-oriented programming is the wrong paradigm for Internet programming!
 - Important: **isolation**, **concurrency**, **asynchronous messages**, **higher-order programming**
 - Unimportant: **inheritance**, **classes**, **methods**, **UML diagrams**, **monitors**


Call Handling Throughput for one CP - AXD 301 release 3.2
Traffic Case: ATM SVC UNI to UNI



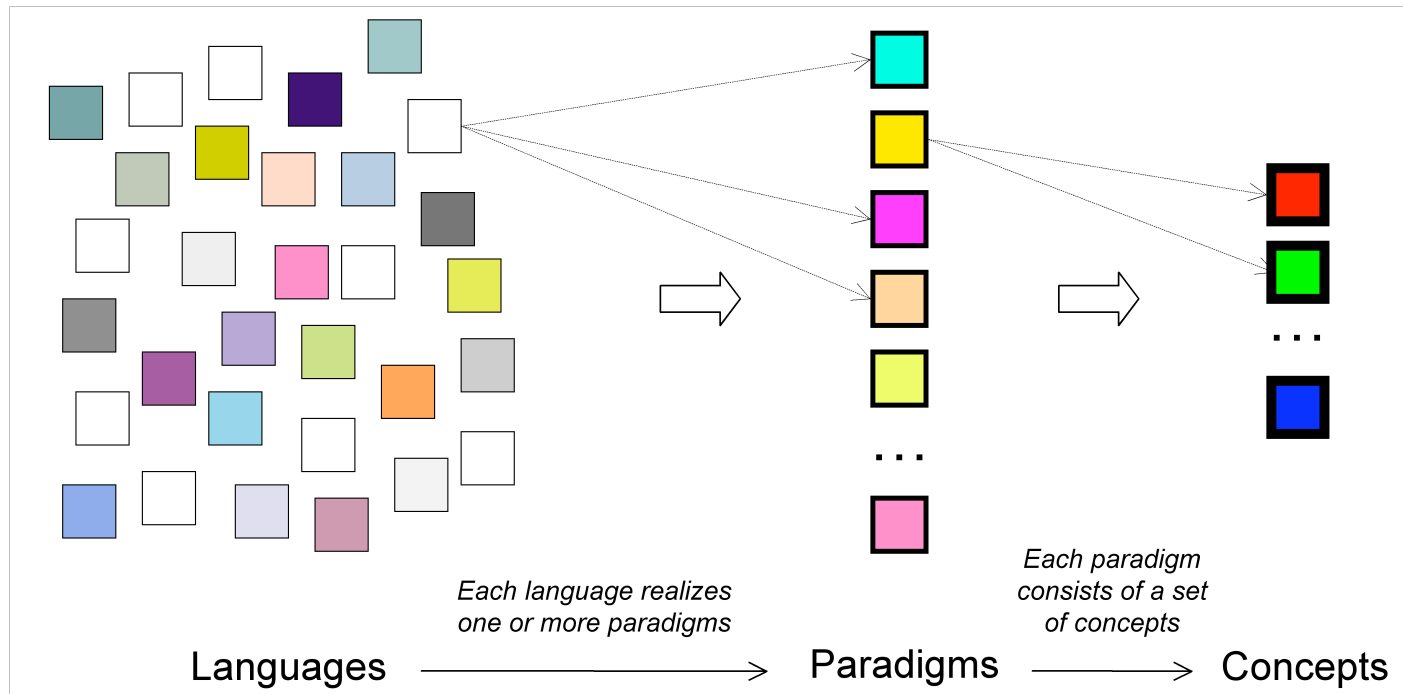
Abbreviations:

CP	Control Processor
ATM	Asynchronous Transfer Mode
SVC	Switched Virtual (ATM) Channel
UNI	User-Network Interface signaling protocol

Paradigms and languages



Paradigms, not languages



- There are lots of programming languages
 - Hundreds are used in industry: you will see very many in your careers!
- There are many fewer programming paradigms and concepts
 - Less than 30 paradigms are considered important
- If you want to understand programming, first understand paradigms!

The creative extension principle



- How do we invent a new paradigm?
- When, in a given paradigm, programs start getting complicated for technical reasons that are unrelated to the problem being solved (e.g., nonlocal program transformations are needed), then there is **a new programming concept waiting to be discovered!**
 - Adding this concept to the paradigm lets programs get simple again
- We saw one example: the concept of **nondeterministic choice**
- Another example is the concept of **exceptions**
 - If the paradigm does not support them (e.g., like in C), then all routines on the call path must test and return error codes
 - If the paradigm supports them (e.g., like in Java), then only the ends of the call path need to be changed (raise and catch exceptions)

Exception handling



Language without exceptions

```
proc {P1 ... E1}
  {P2 ... E2}
  if E2 then ... end
  E1=...
end

proc {P2 ... E2}
  {P3 ... E3}
  if E3 then ... end
  E2=...
end

proc {P3 ... E3}
  {P4 ... E4}
  if E4 then ... end
  E3=...
end

proc {P4 ... E4}
  if (error) then E4=true
  else E4=false end
end
```

Error is handled here

All procedures
on the call path
must be modified

Error appears here

Language with exceptions

```
proc {P1 ...}
  try
    {P2 ...}
  catch E then ... end
end

proc {P2 ...}
  {P3 ...}
end

proc {P3 ...}
  {P4 ...}
end

proc {P4 ...}
  if (error) then
    raise myError end
  end
end
```

Error is handled here

Only the procedures
at the ends
must be modified

Error appears here

Unchanged



So which paradigm is best?

- Each is best for a particular kind of problem the paradigm paradox
 - None is best overall: *“more is not better or worse, only different”*
- The conventional boundaries between paradigms are completely artificial (they exist only for historical reasons)
 - Java is only object-oriented ⇒ **wrong**
 - Scala is functional, object-oriented, and actor-based ⇒ **right**
- A big program almost always needs several paradigms
 - This is why you need to learn about multiple paradigms
- A good language should support several paradigms
 - This is hard for industry languages (Scala and Erlang are moving in the right direction; Java and C++ are stagnating)
 - In your course, you will use Oz: a research language that supports many paradigms (Oz ideas are slowly moving to industry...)