

Tutorial of Oz 2 and the DFKI Oz System

Seif Haridi

SICS, Swedish Institute of Computer Science

Box 1263, SE-164 29 KISTA

Table of Contents

| | |
|---|----|
| Abstract | 4 |
| 1 Basics..... | 5 |
| 1.1 Introduction..... | 5 |
| 1.2 Variables Declaration | 6 |
| 1.3 Primary Oz Types..... | 7 |
| 1.4 Data Types with Structural Equality..... | 8 |
| 1.5 Numbers..... | 9 |
| 1.6 Literals | 9 |
| Lists | 11 |
| 1.7 Virtual Strings | 12 |
| 2 Equality and the Equality Test Operator | 14 |
| 3 Basic Control Structures..... | 17 |
| 3.1 skip..... | 17 |
| 3.2 Conditionals | 17 |
| 3.3 Procedural Abstraction | 19 |
| 3.4 Anonymous Procedures and Variable Initialization..... | 20 |
| 3.5 Pattern Matching | 21 |
| 3.6 Nesting..... | 23 |
| 3.6.1 Functional Nesting..... | 24 |
| 3.7 Procedures as Values..... | 25 |
| 3.8 Control Abstractions..... | 26 |
| 3.9 Exception Handling | 27 |
| 4 Functional Notation..... | 30 |
| 5 Modules and Interfaces..... | 34 |
| 6 Concurrency..... | 37 |
| 6.1 Time..... | 38 |
| 6.2 Stream Communication | 39 |
| 6.3 Thread Priority and Real Time..... | 41 |
| 6.4 Demand-driven Execution | 42 |
| 7 Stateful Data Types | 46 |
| 7.1 Ports..... | 46 |
| 7.2 Chunks | 48 |
| 7.3 Cells | 48 |
| 8 Classes and Objects..... | 52 |
| 8.1 Objects and Classes for Real..... | 53 |
| 8.2 Static Method Calls | 54 |
| 8.3 Inheritance..... | 55 |
| 8.4 Features..... | 57 |
| 8.5 Parameterized Classes..... | 59 |
| 8.6 Self Application..... | 60 |
| 8.7 Attributes..... | 60 |
| 8.8 Private and Protected Methods..... | 61 |
| 8.9 Default Argument Values | 63 |
| 9 Objects and Concurrency..... | 66 |
| 9.1 Locks..... | 66 |
| 9.1.1 Simple Locks..... | 66 |
| 9.1.2 Thread-Reentrant Locks..... | 67 |
| 9.2 Locking Objects | 69 |
| Bibliography | 74 |

Figures

| | |
|---|----|
| • Figure 1 Oz Type Hierarchy | 7 |
| • Figure 2 using a case statement..... | 19 |
| • Figure 3Tree insertion. | 22 |
| • Figure 4 Tree insertion using case statement. | 23 |
| Figure 5 Binary Merge of two lists. | 23 |
| • Figure 6 Binary merge of two lists in nested form..... | 25 |
| • Figure 7. Checking a binary tree. | 25 |
| • Figure 8. Checking a binary tree lazily..... | 26 |
| • Figure 9. The For iterator..... | 26 |
| Figure 10. Checking a binary tree lazily..... | 31 |
| • Figure 11 Checking a binary tree lazily..... | 32 |
| • Figure 12 A concurrent Map function. | 38 |
| • Figure 13 A concurrent Fibonacci function. | 38 |
| • Figure 14 A ‘Ping Pong’ program..... | 39 |
| • Figure 15 Summing the elements in a list. | 39 |
| • Figure 16 Producing volvo's. | 40 |
| • Figure 17 Producing volvo's lazily..... | 43 |
| • Figure 18 Concurrent Composition. | 44 |
| • Figure 19 Concurrent Queue server. | 47 |
| • Figure 20. Implementation of Ports by Cells and Chunks..... | 50 |
| • Figure 21. An Example of Class construction. | 52 |
| • Figure 22. Object Construction. | 53 |
| • Figure 23. Counter Class. | 54 |
| • Figure 24. List Class..... | 55 |
| • Figure 25. Illegal class hierarchy. | 56 |
| • Figure 26. Illegal class hierarchy in method m. | 56 |
| • Figure 27. Parameterized Classes..... | 59 |
| • Figure 28. The class Point. | 61 |
| • Figure 29. The class History Point. | 62 |
| • Figure 30 The class BoundedPoint..... | 64 |
| • Figure 31 The class BHPoint. | 65 |
| • Figure 32 Using Lock. | 68 |
| • Figure 33 An Asynchronous Channel Class | 70 |
| • Figure 34 A Unit Buffer Monitor..... | 72 |
| • Figure 35 Unit Buffer. | 73 |
| • Figure 36 A Bounded Buffer Class. | 73 |

Abstract

This tutorial introduces the programming language Oz 2. Oz 2 is a multiparadigm programming language. It is a high level programming language that is designed for advanced, concurrent, networked, soft real-time, and reactive applications. Oz 2 combines the salient features of object-oriented programming by providing state, abstract data types, classes, objects and inheritance. It provides features of functional programming by providing compositional syntax, first class procedures, and lexical scoping. It provides the salient features of logic and constraint programming such as logic variables, constraints, disjunctions, and programmable search mechanisms. Oz 2 is a concurrent language where users can create dynamically any number of sequential threads. Oz 2 threads are dataflow whose execution proceeds only when data dependencies on variables involved are resolved.

The tutorial covers most of the concepts of Oz 2 in an informal way. It is a suitable as first reading for programmers that want to be able to start quickly writing applications without any particular theoretical background. The document is deliberately informal and, thus complements other DFKI Oz 2 documentations.

1 Basics

1.1 Introduction

A very good starting point is to ask why Oz 2. Well, one rough short answer is that, compared to other existing languages, it is magic! It provides the programmers and system developers with a wide range of programming abstractions to enable them to develop complex applications quickly and robustly. Oz 2 tries to merge several directions of programming language designs into a single coherent one. Yet, Oz 2 is a simple and coherent design. Most of us know the benefits of the various programming paradigms whether object-oriented, functional or constraint logic programming. When we start writing programs in any existing language, we quickly find ourselves confined by the concepts of the underlying paradigm. Oz tries to attack this problem by a coherent design of a language that combines the programming abstractions of various paradigms in clean and simple way.

So, before answering the above question, let us see what Oz 2 is. This is again a difficult question to answer in a few sentences. So, here is the first shot. It is a high level programming language that is designed for modern advanced, concurrent, intelligent, networked, soft real-time, parallel, interactive and pro-active applications. As you see, it is still hard to know what all this jargon means.

- Oz 2 combines the salient features of object-oriented programming, by providing state, abstract data types, classes, objects and inheritance.
- It provides the salient features of functional programming by providing a compositional syntax, first class procedures, and lexical scoping. In fact, every Oz 2 entity is first class, including procedures, threads, classes, methods, and objects.
- It provides the salient features of logic and constraint programming by providing logical variables, disjunctions, flexible search mechanisms and constraint programming.
- It is also a concurrent language where users can create dynamically any number of sequential threads that can interact with each other. However, in contrast to conventional concurrent languages, each Oz 2 thread is a data-flow thread. Executing a statement in Oz 2 proceeds only when all *real* data-flow dependencies on the variables involved are resolved.

Oz 2 has its roots in a paradigm that is known as *concurrent constraint programming*¹. It is a successor to the programming language Oz and its programming system DFKI Oz. This earlier language and system will be called Oz 1 from now on, while the language will be called Oz. There are a number of differences between Oz 1 and Oz 2. The main difference, however, is that Oz 1 is a fine-grained concurrent language where potentially every statement could be executed in its own thread. Oz 2 abandoned this decision and returns to conventional sequential control structure. In spite of this, threads can be created easily and cheaply in Oz 2. Oz 2 is a data-flow language. This feature is retained from Oz 1.

¹ Vijay Saraswat, Concurrent Constraint Programming, MIT press, 1994.

1.2 Variables Declaration

We will initially restrict ourselves to the sequential programming style of Oz. You can think of Oz computations as performed by one sequential process that executes one statement after the other. We call this process a *thread*. This thread has access to a memory called the *store*. It is able to manipulate the store by reading, adding, and updating information. Information is accessed through the notion of *variables*. A thread can access information only through the variables visible to it, directly or indirectly. Oz variables are *single-assignment* variables. In imperative languages like C and Java, a variable can be assigned multiple times. In contrast, single assignment variables can be assigned only once. This notion is known from many languages including data flow, and concurrent logic programming languages. A single assignment variable has a number of phases in its lifetime. Initially it is introduced with unknown value, and later it might be assigned a value, in which case the variable becomes *bound*. Once a variable is bound, it cannot itself be changed. However this does not mean that you cannot model state-change because a variable, as you will see later, could be bound to a cell, which is stateful, i.e. the content of cell can be changed.

A thread executing the statement:

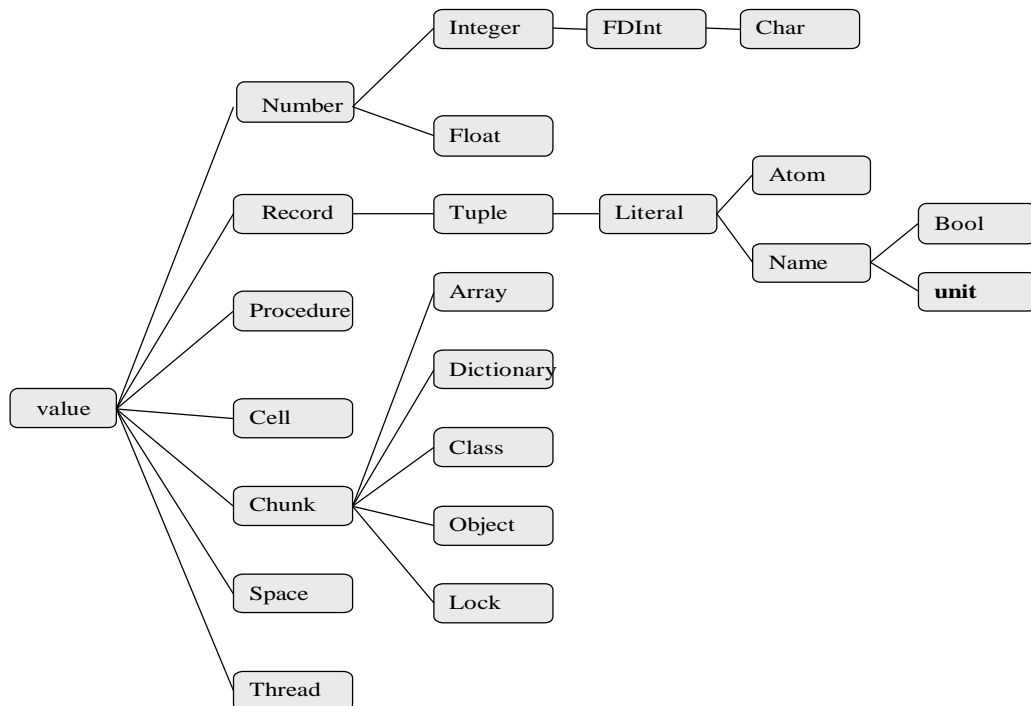
```
local X Y Z in S end
```

Will introduce three single assignment variables *X*, *Y* and *Z* and execute *S* in the scope of these variables. A variable normally starts with an upper-case letter, possibly followed by an arbitrary number of alphanumeric characters. Variables may be also presented textually as any string of printable characters enclosed within back-quote characters, e.g. ``this $ is a variable``. Before the execution of *S* the variables declared, will not have any associated values. We say that the variables are *unbound*. Any variable in an Oz program must be introduced, except for certain pattern matching constructs to be shown later.

Another form of declaration is:

```
declare X Y Z in S
```

This is an open-ended declaration that makes *X*, *Y* and *Z* visible globally in *S*, as well as all statements that follow *S* textually, unless overridden again by another variable declaration of the same textual variables.



• Figure 1 Oz Type Hierarchy

1.3 Primary Oz Types

Oz is a dynamically typed language. Figure 1 shows the type hierarchy of Oz. Any variable, if it ever gets a value, will be bound to a value of one of these types. Most of the types seem familiar to experienced programmers, except probably *Chunk*, *Cell*, *Space*, *FDInt* and *Name*. We will discuss all of these types in due course. For the impatient reader here are some hints. The *Chunk* data type allows users to introduce new abstract data types. *Cell* introduces the primitive notion of state-container and state modification. *Space* will be needed for advanced problem solving using search techniques. *FDInt* is the type of finite domain that is used frequently in constraint programming, constraint satisfaction, and operational research. *Name* introduces anonymous unique unforgeable tokens.

The language is dynamically typed in the sense that when a variable is introduced its type, as well as its value, is unknown. Only when the variable is bound to an Oz value, its type becomes determined.

Hello World

Let us do like everybody else. If you are unfamiliar with Oz, here is your first Oz program. Start your oz system. You typically start it by writing the command: `oz` and return. This will start the system having the Emacs editor as your interface. You will see two buffers. The upper buffer is called `Oz`, where you can enter programs and the lower buffer is called `*Oz Compiler*` where you can see the result of compiling Oz

programs. There is also a third buffer `*Oz Emulator*` showing the status of the emulator. This buffer is your standard input, and standard output. You may switch between the compiler and the emulator buffer through the Oz pull-down menu that appears in your EMACS menu bar. Use the entry `Show/Hide` to switch between the compiler and the emulator buffer.

If you feed the program show below:

```
{Show 'Hello World'}
```

It will print the string `'Hello World'` in your standard output. In fact, the main reason for showing this program is to get you accustomed with one of the unconventional syntactic aspects of Oz, namely the syntax of *procedure calls/applications*. `{Show 'Hello World'}` is a procedure application of `Show` on the single atom argument `'Hello World'`. `Show` is actually a pre-declared global variable in your initial environment that got bound to the printing procedure when the `System` module was loaded at the start of the Oz system. Procedure application in Oz syntactically follows other functional languages, e.g. SCHEME, with the exception of using braces instead of parentheses.

The DFKI Oz system provides a number of interesting tools that are accessible through the Oz menu, or can be called from your Oz Program. Instead of using `Show` try `Browse`.

You can learn about the DFKI Oz programming system, its user interface and associated tools by looking to 'The DFKI Oz User's Manual', *The Oz documentation series*.

The module `System` is defined in 'The DFKI Oz User's Manual'.

Adding Information

In Oz, there are few ways of adding information to the store or (said differently) of binding a variable to a value. The most common form is using the *equality* infix operator `=`. For example, given that the variable `x` is declared the following statement:

```
x = 1
```

will bind the unbound variable `x` to the integer 1, and add this information to the store. Now, if `x` is already assigned the value 1, the operation is considered as performing a test on `x`. If `x` is already bound to an incompatible value, i.e. to any other value different from 1, a proper *exception* will be raised. Exception handling is described later.

1.4 Data Types with Structural Equality

The hierarchy starting from `Number` and `Record` in Figure 1 defines the data types of Oz whose members (values) are equal only if they are structurally similar. For example two numbers are equal if they have the same type, or one is a subtype of the other, and have the same value. For example, both are integers and are the same number, or both are lists and their head elements are equal as well as their respective tail lists. Structural equality allows values to be equivalent even if they are replicas occupying different physical memory location.

1.5 Numbers

The following program, introduces three variables `I`, `F` and `C`. It assigns `I` an integer, `F` a float, and `C` the character 't' in this order. It then displays the list consisting of `I`, `F`, and `C`.

```
local I F C in
  I = 5
  F = 5.5
  C = &t
  {Browse [I F C]}
end
```

Oz support binary, octal, decimal and hexadecimal notation for integers, which can be arbitrary large. An octal starts with a leading 0, and a hexadecimal starts with a leading 0x or 0X. Floats are different from integers and must have decimal points. Other examples of floats are shown where '~' is unary minus:

```
~3.141 4.5E3 ~12.0e~2
```

In Oz, there is no automatic type conversion, so `5.0 = 5` will raise an exception. Of course, there are primitive procedures for explicit type conversion. These and many others can be found in [7] Characters are a subtype of integers in the range of 0, . . . , 255. The standard ISO 8859-1 coding is used (not unicode). Printable characters have external representation, e.g. `&0` is actually the integer 48, and `&a` is 97. Some control characters have also a representation e.g. `&\n` is a new line. All characters can be written as `&\ooo`, where `o` is an octal digit.

Operations on characters, integers, and floats can be found in the library modules `Char`, `Float`, and `Int`. Additional generic operations on all numbers are found in the module `Number`.

1.6 Literals

Another important category of atomic types, i.e. types whose members have no internal structure, is the category of literals. Literals are divided into atoms and names. An Atom is symbolic entity that has an identity made up of a sequence of alphanumeric characters starting with a lower case letter, or arbitrary printable characters enclosed in quotes. For example:

```
a foo '=' ':= ' 'OZ 2.0' 'Hello World'
```

Atoms have an ordering based on lexicographic ordering.

Another category of elementary entities is Name. The only way to create a name is by calling the procedure `{NewName X}` where `X` is assigned a new name that is guaranteed to be worldwide unique. Names cannot be forged or printed. As will be seen later, names play an important role in the security of Oz programs. A subtype of Name is `Bool`, which consists of two names protected from being redefined by having the reserved keywords `true` and `false`. Thus a user program cannot redefine them, and mess up all programs relying on their definition. There is also the type `Unit` that consists of the single name `unit`. This is used as synchronization token in many concurrent programs.

```

local X Y B in
  X = foo
  {NewName Y}
  B = true
  {Browse [X Y B]}
end

```

Records and Tuples

Records are structured compound entities. A record has a *label* and a fixed number of components or arguments. There are also records with a variable number of arguments that are called *open records*. For now, we restrict ourselves to ‘closed’ records. The following is a record:

```
tree(key: I value: Y left: LT right: RT)
```

It has four arguments, and the label *tree*. Each argument consists of a pair *Feature:Field*, so the features of the above record is *key*, *value*, *left*, and *right*. The corresponding fields are the variables *I*, *Y*, *LT*, and *RT*. It is possible to omit the features of a record reducing it to what is known from logic programming as a compound term. In Oz, this is called a *tuple*. So, the following tuple has the same label and fields as the above record:

```
tree(I Y LT RT)
```

It is just a syntactic notation for the record:

```
tree(1:I 2:Y 3:LT 4:RT)
```

where the features are integers starting from 1 up to the number of fields in the tuple. The following program will display a list consisting of two elements one is a record, and the other is tuple having the same label and fields:

```

declare T I Y LT RT W in
  T = tree(key:I value:Y left:LT right:RT)
  I = seif
  Y = 43
  LT = nil
  RT = nil
  W = tree(I Y LT RT)
  {Browse [T W]}

```

The display will show:

```

[tree(key:seif value:43 left:nil right:nil)
 tree(seif 43 nil nil)]

```

Operations on records

We discuss some basic operations on records. Most operations are found in the module *Record*. To select a field of a record component, we use the infix operator ‘.’

Record.Feature

```

% Selecting a Component
{Browse T.key}
{Browse W.1}
% will show seif twice on the display.
seif
seif

```

The *arity* of a record is a list of the features of the record sorted lexicographically. To display the arity of a record we use the procedure *Arity*. The procedure application

{Arity R X} will execute once R is assigned a record, and will bind x to the arity of the record. Executing the following statements

```
% Getting the Arity of a Record
local X in {Arity T X} {Browse X} end
local X in {Arity W X} {Browse X} end
```

will display

```
[key left right value]
[1 2 3 4]
```

Another useful operation is conditionally selecting a field of a record. The operation `CondSelect` takes a record R, a feature F, and a default field-value D, and a result argument X. If the feature F exists in R, R.F is assigned to X, otherwise the default value D is assigned to X. `CondSelect` is not really a primitive operation. It is definable in Oz. The following statements:

```
% Selecting a component conditionally
local X in {CondSelect W key eevee X} {Browse X} end
local X in {CondSelect T key eevee X} {Browse X} end
```

will display

```
eevee
seif
```

A common infix tuple-operator used in Oz is `'#'`. So, `1#2` is a tuple of two elements, and observe that `1#2#3` is a single tuple of three elements:

```
'#'(1 2 3)
```

and not `1#(2#3)`. With the `'#'` operator there is no empty tuple notation, and single element tuple is written as `'#'(X)`.

The operation `{AdjoinList R LP S}` takes a record R a list of feature-field pairs, and returns in S a new record such that:

- The label of R is equal to the label of S.
- S has the components that are specified in LP in addition to all components in R that do not have a feature occurring in LP.

```
local S in
  {AdjoinList tree(a:1 b:2) [a#3 c#4] S}
  {Show S}
end
% gives S=tree(a:3 b:2 c:4)
```

Lists

As in many other symbolic programming languages, e.g. Scheme and Prolog, *lists* forms an important class of data structures in Oz. The category of lists does not belong to a single data type in Oz. They are rather conceptual structure. A list is either the atom `nil` representing the empty list, or is a tuple using the infix operator `'|'` and two arguments which are respectively the head and the tail of the list. Thus, a list of the first three natural numbers is represented as:

```
1|2|3|nil
```

Another convenient special notation for a *closed list*, i.e. list with determined number of elements is:

```
[1 2 3]
```

The above notation is used only for closed list, so a list whose first two elements are `a` and `b`, but whose tail is the variable `x` looks like:

```
1 | 2 | x
```

One can also use the standard notation for lists:

```
' | '(1 ' | '(2 x))
```

Further notational variant is allowed for lists whose elements correspond to character codes. Lists written in this notation are called *strings*, e.g.

```
"OZ 2.0"
```

is the list

```
[79 90 32 50 46 48]
```

1.7 Virtual Strings

A virtual string is special a tuple that represents a string with virtual concatenation, i.e. the concatenation is performed when really needed. Virtual strings are used for I/O with files, sockets, and windows. All atoms, except `nil` and `#`, as well as numbers, strings, or `#`-labeled tuples can be used to compose virtual strings. Here is one example:

```
123#"- "#23#" is "#100
```

represents the string

```
"123-23 is 100"
```

For each data type discussed in section, there is a corresponding module in the DFKI Oz systems. The modules define operations on the corresponding data type. You can find more about these operation in 'The Oz Standard Modules', *The Oz documentation series*.

2 Equality and the Equality Test Operator

We have so far shown simple examples of the equality statement, e.g.

```
W = tree(I Y LT LR)
```

These were simple enough to understand intuitively what is going on. However, what happens when two unbound variables are equated $X = Y$, or when two large data structures are equated. Here is a short explanation. We may think of the store as a dynamically expanding array of memory cells. Each cell is labeled by a single-assignment variable. When a variable X is introduced a cell is created in the store, labeled by X , and its value is *unknown*. At this point, the cell does not possess any value; it is empty as a container that might be filled later.

A variable with an empty cell is an *unbound* variable. The cell is flexible enough to contain any arbitrary Oz value. The operation

```
W = tree(1:I 2:Y 3:LT 4:LR)
```

stores the record structure in the cell associated with W . Notice that we are just getting a graph structure. The cell contains a record with four fields. The fields contain arcs pointing to the cells labeled by I , Y , LT , and LR respectively. Each arc in turn is labeled by the corresponding feature of the record. Given two variables X and Y , $X = Y$ will try to *merge* their respective cells. Now we are in a position to give a reasonable account for the merge operation of $X = Y$, known as the *incremental tell* operation.

- X and Y label the same cell: the operation is completed successfully.
- $X(Y)$ is unbound: merge the cell of $X(Y)$ with the cell of $Y(X)$. Merging is done by making the cell of X point to that of Y ; X and Y are now considered to be labeling the same cell. Conceptually the original cell of X has been discarded.

X and Y are labels of different cells containing the records R_x and R_y respectively:

- R_x and R_y have different labels, arities, or both: the operation is completed, and an exception is raised.
- Otherwise, the arguments of R_x and R_y with the same feature are pair-wise merged in arbitrary order.

In general the two graphs, to be merged, could have cycles. However any correct implementation of the merge operation will remember cell pairs for which an attempt to merge has been made earlier, can consider the operation to be successfully performed.

When a variable is no longer accessible, the cell and its associated variable is reclaimed by a process known as garbage collection.

Here are some examples of successful equality operations:

```
local X Y Z in
  f(1:X 2:b) = f(a Y)
  f(Z a) = Z
  {Browse [X Y Z]}
end
```

will show `[a b R14=f(R14 a)]` in the browser. `R14=f(R14 a)` is the external representation of a cyclic graph.

To be able to see the finite representation of *Z*, you have to switch the browser to **Minimal Graph** presentation mode. Choose the Option menu, Representation field, and click on Minimal Graph.

The Browser is described in 'The DFKI Oz User's Manual'.

The following example shows, what happens when variables with incompatible values are equated.

```
declare X Y Z in
  X = f(c a)
  Y = f(Z b)
  X = Y
```

The incremental tell of *X = Y* will assign *Z* the value *c*, but will also raise an exception that is caught by the system, when it tries to equate *a* and *b*.

equality test operator ==

The basic procedure `{`==` X Y R}` tries to test whether *X* and *Y* are equal or not, and returns the result in *R*.

- It returns the Boolean value **true**, if the graphs starting from the cells of *X* and *Y* have the same structure, with each pair-wise corresponding cells having identical Oz values or are the same cell.
- It returns the Boolean value **false**, if the graphs have different structure, or some pair-wise corresponding cells have different values.
- It suspends when it arrives at pair-wise corresponding cells that are different, but at least one of them is unbound.

Now remember if a procedure suspends, the whole thread suspends! This does not seem useful however, as you will see later, it becomes a very useful operation when multiple thread start interacting with each other.

The equality test is normally used as a functional expression, rather than a statement. As shown is the following example:

```
% See, lists are just tuples, which are just records
local L1 L2 L3 Head Tail in
  L1 = Head|Tail
  Head = 1
  Tail = 2|nil

  L2 = [1 2]
  {Browse L1==L2}

  L3 = '|'(1:1 2:'|(2 nil))
  {Browse L1==L3}
end
```


3 Basic Control Structures

We have already seen basic statements in Oz. Introducing new variables and sequencing of statements:

$$S_1 S_2$$

Reiterating again, a thread executes statements in a sequential order. However a thread, contrary to conventional languages, may suspend in some statement, so above a thread has to complete execution of S_1 , before starting S_2 . In fact, S_2 may not be executed at all, if an exception is raised in S_1 .

3.1 skip

The statement **skip** is the empty statement.

3.2 Conditionals

Simple Conditionals. A statement having the following form:

if $X_1 \cdots X_n$ **in** C **then** S_1 **else** S_2 **end**

where C is a sequence of simple equalities, is called a simple conditional. The subexpression **if** $X_1 \cdots X_N$ **in** C **then** S_1 is usually called a *clause*, and **if** $X_1 \cdots X_N$ **in** C is the *condition* of the clause. A thread executing such as conditional will first check whether the condition: "There are $X_1 \cdots X_N$ such that C is true" is satisfied or falsified by the current state of the store. If the condition is satisfied, statement S_1 is executed; if it is falsified, the thread executes S_2 ; and if neither hold, the thread suspends.

A simple equality have in general the form $x = t$ where t is a record or a variable. Examples of simple conditions follow.

```
X = true           % X is bound to true in the store
Y Z in X = f(Y Z)  % X is bound to a tuple of 2 arguments
X = Y              % X and Y label the same cell in the store
```

Comparison Procedures. Oz provides a number of built-in tertiary procedures used for comparison. These include `==` that we have seen earlier as well as `\=`, `=<`, `<`, `>=`, `>`, **andthen**, and **orelse**. Common to these procedures is that they are used as Boolean functions in an infix notation. The following example illustrates the use of a conditional in conjunction with the greater-than operator `>`. In this example z is bound to the maximum of x and y , i.e. to y :

```

local X Y F Z in
  X = 5
  Y = 10
  F = X > Y
  if F = true then Z = X
  else Z = Y end
end

```

Parallel Conditional. A parallel conditional is of the form:

```

if C1 then S1
[] C2 then S2
  ⋮
else SN end

```

A parallel conditional is executed by evaluating all conditions $C_1 \cdots C_{N-1}$ in an *arbitrary* order, possibly concurrently. If one of the conditions say C_i is true, its corresponding statement S_i is chosen. If all conditions are false, the else statement S_N is chosen, otherwise the executing thread suspends. Parallel conditionals are useful mostly in concurrent programming, e.g. for programming time-out on certain events. However, it is often used in a deterministic manner with mutually exclusive conditions. So, the above example could be written as:

```

local X Y F Z in
  X = 5
  Y = 10
  F = X >= Y
  if F = true then
    Z = X
  [] F = false then
    Z = Y
  end
end

```

Notice that the **else** part is missing. This is just a shorthand notation for an **else** clause that raises an exception.

Case Statement. Oz provides an alternative conditional syntax that encourages the use of the simple form of conditions. This is called the *case* statement. The following is the simplest form:

```

case B then S1 else S2 end

```

where B is a simple condition that should be evaluated to a Boolean value. If B is **true** S_1 is executed, otherwise if B is **false** S_2 is executed. This is equivalent to:

```

if B = true then S1 else S2 end

```

Our example using if-statement could be now written as shown below.

```

local X Y Z in
  X = 5 Y = 10
  case X >= Y then Z = X else Z = Y end
end

```

- Figure 2 using a case statement.

Since a case-statement is defined in terms of an if-statement, it is merely a notational convenience.

3.3 Procedural Abstraction

Procedure definitions. Procedure definition is a primary abstraction in Oz. A procedure can be defined, passed around as argument to another procedure, or stored in a record. A procedure definition is a statement that has the following structure.

```

proc {P X1...XN} S end

```

Assume that the variable P is already introduced; executing the above statement will create a procedure, consisting of a unique name α and an abstraction $I(X_1 \cdots X_N) \cdot S$. This pair is stored in the memory cell labeled by P . A procedure in Oz has a unique identity, given by its name, and is distinct from all other procedures. Two procedure definitions are always different, even if they look similar. Procedures are the first Oz values that we encounter, whose equality is based on name equality. Others include threads, cells, and chunks.

On Lexical Scoping. In general, the statement S in a procedure definition will have many ‘syntactic’ variable occurrences. A syntactic variable is called an identifier to distinguish it from the single assignment variable that is created at runtime. Some identifier occurrences in S are *syntactically bound* while others are *free*. An identifier occurrence X in S is bound if it is in the scope of the procedure formal-parameter X , or is in the scope of a variable introduction statement that introduces X ². Otherwise, the identifier occurrence is free. Each free identifier occurrence in a program is eventually bound by the closest textually surrounding identifier-binding construct.

We have already seen how to apply (call) a procedure. Let us now show our first procedure definition. In Figure 2, we have seen how to compute the maximum of two numbers or literals. We abstract this code into a procedure.

² This rule is approximate, since class methods and patterns bind identifier occurrences.

```

local Max X Y Z in
  proc {Max X Y Z}
    case X >= Y then
      Z = X
    else Z = Y end
  end
  X = 5 Y = 10
  {Max X Y Z} {Browse Z}
end

```

3.4 Anonymous Procedures and Variable Initialization

One could ask why is a variable bound to a procedure in a way that is different from the variable being bound to record: $X = Value$? The answer is that what you see is just a syntactic variant of the equivalent form

$$P = \mathbf{proc} \{ \$ X_1 \cdots X_N \} S \mathbf{end}$$

where the R.H.S. defines an *anonymous procedural value*. This is equivalent to

$$\mathbf{proc} \{ P X_1 \cdots X_N \} S \mathbf{end}$$

In Oz, we can initialize a variable immediately while it is being introduced by using a *variable-initialization equality*

$$X = Value \text{ or } R = Value \text{ (} X \text{ occurs in the record } R \text{)}$$

between **local** and **in**, in the statement **local** ... **in** ... **end**. So the previous example could be written as follows, where we also use anonymous procedures.

```

local
  Max = proc { $ X Y Z }
    case X >= Y then Z = X
    else Z = Y end
  end
  X = 5
  Y = 10
  Z
in
  {Max X Y Z} {Browse Z}
end

```

Now let us understand variable initialization in more detail. The general rule says that: in a variable-initialization equality, only the variables occurring on the L.H.S. of the equality are the ones being introduced. Consider the following example:

```

local
  Y = 1
in
  local
    M = f(M Y)
    [X1 Y] = L
    L = [1 2]
  in {Browse [M L]} end
end

```

First *Y* is introduced and initialized in the outer **local ... in ... end**. Then, in the inner **local ... in ... end** all variables on the L.H.S. are introduced, i.e. *M*, *Y*, *X1*, and *L*. Therefore the outer variable *Y* is invisible in the innermost **local ... end** statement. The above statement is equivalent to:

```

local Y in
  Y = 1
  local M X1 Y L in
    M = f(M Y)
    L = [X1 Y]
    L = [1 2]
    {Browse [M L]}
  end
end

```

If we want *Y* to denote the variable in the outer scope, we have to suppress the introduction of the inner *Y* in the L.H.S. of the initializing equality by using an exclamation mark ‘!’ as follows. An exclamation mark ‘!’ is only meaningful in the L.H.S. of an initializing equality.

```

local
  Y = 1
in
  local
    M = f(M Y)
    [X1 !Y] = L
    L = [1 2]
  in {Browse [M L]}
  end
end

```

3.5 Pattern Matching

Let us consider a very simple example: insertion of elements in a binary tree. A binary tree is either empty, represented by *nil*, or is a tuple of the form *tree(Key Value TreeL TreeR)*, where *Key* is a key of the node with the corresponding value *Value*, and *TreeL* is the left subtree having keys less than *Key*, and *TreeR* is the right subtree having keys greater than *key*. The procedure *Insert* takes four arguments, three of them are input arguments *Key*, *Value* and *TreeIn*, and one output argument *TreeOut* to be bound to the resulting tree after insertion.

The program is shown in Figure 3. The symbol ‘?’ before *TreeOut* is a voluntary *documentation comment* denoting that the argument plays the role of an output argument. The procedure works by cases as obvious. First depending on whether the tree is empty or not, and in the latter case depending on a comparison between the key of the node in the tree and the input key. Notice the use of **case ... then ... elsecase ... else ... end** with the obvious meaning.

```

declare
proc {Insert Key Value TreeIn ?TreeOut}
  if TreeIn = nil then TreeOut = tree(Key Value nil nil)
  [] K1 V1 T1 T2 in TreeIn = tree(K1 V1 T1 T2) then
    case Key == K1 then TreeOut = tree(Key Value T1 T2)
    elsecase Key < K1 then
      local T in
        TreeOut = tree(K1 V1 T T2)
        {Insert Key Value T1 T}
      end
    else
      local T in
        TreeOut = tree(K1 V1 T1 T)
        {Insert Key Value T2 T}
      end
    end
  end
end

```

- Figure 3 Tree insertion.

In Figure 3, it is tedious to introduce the local variables in the clause

```

[] K1 V1 T1 T2 in TreeIn = tree(K1 V1 T1 T2) then ...

```

Oz provides a pattern-matching **case** statement, which allows implicit introduction of variables in the patterns. Two forms exist for the case-statement:

```

case E
of Pattern1 then S1
elseif Pattern2 then S2
elseif ...
else S end

```

and

```

case E
of Pattern1 then S1
[] Pattern2 then S2
[] ...
else S end

```

All variables introduced in *Pattern_i* are implicitly declared, and have a scope stretching over the corresponding clause. In the first case-matching statement, the patterns are tested in order. In the second, called *parallel* case-matching statement, the order is indeterminate. The **else** part may be omitted, in which case an exception is raised if all matches fail. Again, in each pattern one may suppress the introduction of new local variable by using '!'. For example, in the following example:

```

case f(X1 X2) of f(!Y Z) then ... else ... end

```

X1 is matched against the value of the external variable Y. Now remember again that the case statement and its executing thread may suspend if X1 is insufficiently instantiated to decide the result of the matching. Have all this said, Figure 4 shows the tree-insertion procedure using a matching case-statement. We have also reduced the syntactic nesting by abbreviating:

```

    local T in
        TreeOut = tree( ...T... )
        {Insert ... T}
    end
into:

    T in
        TreeOut = tree( ... T ... )
        {Insert ... T}

```

```

% case for pattern matching
proc {Insert Key Value TreeIn ?TreeOut}
    case TreeIn
    of nil then TreeOut = tree(Key Value nil nil)
    [] tree(K1 V1 T1 T2) then
        case Key == K1 then TreeOut = tree(Key Value T1 T2)
        elseif Key < K1 then T in
            TreeOut = tree(K1 V1 T T2)
            {Insert Key Value T1 T}
        else T in
            TreeOut = tree(K1 V1 T1 T)
            {Insert Key Value T2 T}
        end
    end
end
end

```

- Figure 4 Tree insertion using case statement.

The expression E we may match against, could be any record structure, and not just a variable. This allows multiple argument matching, as shown in Figure 5, which depicts a classical nondeterministic stream-merge procedure. Here the use of parallel-case statement is essential.

```

proc {Merge Xs Ys ?Zs}
    case Xs#Ys
    of nil#Ys then Zs = Ys
    [] Xs#nil then Zs = Xs
    [] (X|Xr)#Ys then Zr in
        Zs = X|Zr
        {Merge Ys Xr Zr}
    [] Xs#(Y|Yr) then Zr in
        Zs = Y|Zr
        {Merge Yr Xs Zr}
    end
end
end

```

- Figure 5 Binary Merge of two lists.
-

3.6 Nesting

Let us use our `Insert` procedure as defined in Figure 4. The following statement inserts a few nodes in an initially empty tree. Note that we had to introduce a number of intermediate variables to perform our sequence of procedure calls.

```

local T0 T1 T2 T3 in
  {Insert seif 43 nil T0}
  {Insert eevee 45 T0 T1}
  {Insert rebecca 20 T1 T2}
  {Insert alex 17 T2 T3}
  {Browse T3}
end

```

Oz provides syntactic support for nesting one procedure call inside another statement at an expression position. So, in general:

```

local Y in
  {P ... Y ...}
  {Q Y ... }
end

```

could be written as:

```

{Q {P ... $ ...} ... }

```

using ‘\$’ as a *nesting marker*, and thereby the variable Y is eliminated. The rule, to revert to the flattened syntax is that, a nested procedure call, inside a procedure call, is moved *before* the current statement; and a new variable is introduced with one occurrence replacing the nested procedure call, and the other occurrence replacing the nesting marker.

3.6.1 Functional Nesting

Another form of nesting is called functional nesting: a procedure {P X ... R} could be considered as a function; its result is the argument R. Therefore

{P X ...} could be considered as a function call that can be inserted in any expression instead of the result argument R. So {Q {P X ...} ...} is equivalent to:

```

local R in
  {P X ... R}
  {Q R ... }
end

```

Now back to our example, a more concise form using functional nesting is:

```

{Browse {Insert alex 17
          {Insert rebecca 20
            {Insert eevee 45 {Insert seif 43
              nil}}}}}

```

There is one more rule to remember. It has to do with a nested application inside a record or a tuple as in:

```

Zs = X | {Merge Xr Ys $}

```

Here, the nested application goes *after* the record construction statement; do you see why? Therefore, we get

```

local Zr in
  Zs = X | Zr
  {Merge Xr Ys Zr}
end

```

We can now rewrite our Merge procedure as shown in Figure 5, where we use nested application.


```

proc {Merge Xs Ys ?Zs}
  case Xs#Ys
  of nil#Ys then Zs = Ys
  [] Xs#nil then Zs = Xs
  [] (X|Xr)#Ys then Zs = X|{Merge Ys Xr $}
  [] Xs#(Y|Yr) then Zs = Y|{Merge Yr Xs $}
  end
end

```

- Figure 6 Binary merge of two lists in nested form.

3.7 Procedures as Values

Since we have been inserting elements in binary trees, let us define a program that checks if a data structure is actually a binary tree. The procedure `BinaryTree` shown in Figure 7 checks a structure to verify whether it is a binary tree or not, and accordingly returns **true** or **false** in its result argument B.

Notice that we also defined the auxiliary local procedure `And`.

```

% What is a binary tree?
local
  proc {And B1 B2 ?B}
    case B1 then
    case B2 then B = true else B = false end
    else B = false end
    end
  in proc {BinaryTree T ?B}
    case T
    of nil then B = true
    [] tree(K V T1 T2) then
      {And {BinaryTree T1} {BinaryTree T2} B}
    else B = false end
    end
  end
end

```

- Figure 7. Checking a binary tree.

Consider the call `{And {BinaryTree T1} {BinaryTree T2} B}`. It is certainly doing unnecessary work. According to our nesting rules, it evaluates its second argument even if the first is **false**. One can fix this problem by making a new procedure `AndThen` that takes as its first two arguments two procedures, and calls the second procedure only if the first returns **false**; thus, getting the effect of delaying the evaluation of its arguments until really needed. The procedure is shown Figure 8. `AndThen` is the first example of a *higher-order procedure*, i.e. a procedure that takes other procedures as arguments, and may return other procedures as results. In our case, `AndThen` just returns a Boolean value. However, in general, we are going to see other examples where procedures return procedures as result. As in functional languages, higher order procedures are invaluable abstraction devices that help creating generic reusable components.

```

local
  proc {AndThen BP1 BP2 ?B}
    case {BP1} then
      case {BP2} then B = true else B = false end
    else B = false end
  end
in proc {BinaryTree T ?B}
  case T
  of nil then B = true
  [] tree(K V T1 T2) then
    {AndThen proc {$ B1}{BinaryTree T1 B1} end
              proc {$ B2}{BinaryTree T2 B2} end
              B}
  else B = false end
end
end

```

- Figure 8. Checking a binary tree lazily.

3.8 Control Abstractions

Higher-order procedures are used in Oz to define various control abstractions. In modules `Control` and `List` as well as many others, you will find many control abstractions. Here are some examples. The procedure `{For From To Step P}` is an iterator abstraction that applies the unary procedure `P` (normally saying the procedure `P/1` instead) to integers from `From` to `To` proceeding in steps `Step`. Notice that use of the empty statement `skip`. Executing `{For 1 10 1 Browse}` will display the integers `1 2 ... 10`.

```

local
  proc {HelpPlus C To Step P}
    case C=<To then {P C} {HelpPlus C+Step To Step P}
    else skip end
  end
  proc {HelpMinus C To Step P}
    case C>=To then {P C} {HelpMinus C+Step To Step P}
    else skip end
  end
in proc {For From To Step P}
  case Step>0 then {HelpPlus From To Step P}
  else {HelpMinus From To Step P} end
end
end

```

- Figure 9. The For iterator.

Another control abstraction that is often used is the `ForAll/2` iterator defined in the `List` module. `ForAll/2` applies a unary procedure on all the elements of a list, in the order defined by the list. Think what happens if the list is produced incrementally by another concurrent thread?

```

proc {ForAll Xs P}
  case Xs
  of nil then skip
  [] X|Xr then
    {P X}
    {ForAll Xr P}
  end
end

```

3.9 Exception Handling

Oz incorporates an exception handling mechanism that allows safeguarding programs against exceptional and/or unforeseeable situations at run-time. It is also possible to raise and handle user-defined exceptions.

An exception is any expression E . To raise the exception E , one executes the following statement:

```
raise E end
```

Here is a simple example:

```

proc {Eval E ?R}
  case E
  of plus(X Y) then {Browse X+Y}
  [] times(X Y) then {Browse X*Y}
  else raise illFormedExpression(E) end
end
end

```

The basic exception handling statement is called a try-statement, and has the following form:

```

try S catch
  Pattern1 then S1
  [] Pattern2 then S2
  ⋮
  [] PatternN then SN
end } optional

```

Execution of this statement is equivalent to executing S if S does not raise an exception. If S raises exception E and E matches one of the patterns $Pattern_i$, control is passed to the corresponding statement S_i . If E does not match any pattern the exception is propagated outside the try-statement until eventually caught by the system, which catches all escaped exceptions.

```

try
  {ForAll [plus(5 10) times(6 11) min(7 10)] Eval}
catch
  illFormedExpression(X) then {Browse '** X **'}
end

```

A try-statement may also specify a final statement S_{FINAL} , which is executed on normal as well as on exceptional exit.

```

try S catch
  Pattern1 then S1
  [] Pattern2 then S2
  ⋮
  [] PatternN then SN } optional
finally SFINAL end

```

Assume that 'F'³ is an opened file; the procedure `Process/1` manipulates the file in some way; and the procedure `CloseFile/1` closes the file. The following program ensures that the F is closed upon normal or exceptional exit.

```

try
  {Process F}
catch
  illFormedExpression(X) then {Browse '** X **'}
finally {CloseFile F} end

```

System Exceptions

The exceptions raised by the Oz system are records with one of the labels: `failure`, `error`, and `system`.

- `failure`: indicates the attempt to perform an inconsistent equality operation on the store of Oz, recall section 2.
- `error`: indicates a runtime error which should not occur.
- `system`: indicates a runtime condition, i.e., an unforeseeable situation like a closed file or window.

The exact format of Oz system-exceptions is in an experimental state and therefore remains the user is only advised to rely of the label only as in the following example:

```

try 1=2
catch
  failure(...) then {Show caughtFailure}
end

```

Here the pattern `failure(...)` catches any record whose label is `failure`. When an exception is raised but not handled, an error message is printed in the emulator window (standard output).

³ We will see how input/output is handled later.

4 Functional Notation

Oz provides functional notation as syntactic convenience. We have seen that a procedure call:

$$\{P X_1 \dots X_N R\}$$

could be used in a nested expression as a function call:

$$\{P X_1 \dots X_N\}$$

Oz also allows functional abstractions directly as syntactic notation for procedures. Therefore, the following function definition:

fun $\{F X_1 \dots X_N\} S E$ **end**

where S is a statement and E is an expression corresponds to the following procedure definition:

proc $\{F X_1 \dots X_N R\} S R = E$ **end**

The exact syntax for functions as well as their unfolding into procedure definitions is defined in *Oz Notation, the Oz Documentation Series*.

Here we rely on the reader's intuition. Roughly speaking, the general rule for syntax formation of functions looks very similar to how procedures are formed. With the exception that, whenever a thread of control in a procedure ends in a statement, the corresponding function ends in an expression.

The program shown in Figure 10 is the functional equivalent to the program shown in Figure 8. Notice how the function `AndThen/2` is unfolded into the procedure `AndThen/3`. Below we show a number of steps that give some intuition of the transformation process. All the intermediate forms are legal Oz programs.

```
fun {AndThen BP1 BP2}
  case {BP1} then
    case {BP2} then true else false end
  else false end
end
```

Make a procedure by introducing a result variable B:

```
proc {AndThen BP1 BP2 B}
  B = case {BP1} then
    case {BP2} then true else false end
  else false end
end
```

Move the result variable into the outer *case-expression* to make it a *case-statement*:

```

proc {AndThen BP1 BP2 B}
  case {BP1} then
    B = case {BP2} then true else false end
  else B = false end
end

```

Move the result variable into the inner *case-expression* to make it a *case-statement*, and we are done:

```

proc {AndThen BP1 BP2 B}
  case {BP1} then
    case {BP2} then B = true else B = false end
  else B = false end
end
% Syntax Convenience: functional notation
local
  fun {AndThen BP1 BP2}
    case {BP1} then
      case {BP2} then true else false end
    else false end
  end
  fun {BinaryTree T}
    case T
    of nil then true
    [] tree(K V T1 T2) then
      {AndThen fun {$}{BinaryTree T1} end
              fun {$}{BinaryTree T2} end}
    else false end
  end
end
end

```

- Figure 10. Checking a binary tree lazily.

If you are a functional programmer, you can cheer up! You have your functions, including higher-order ones, and similar to lazy functional languages Oz allows certain forms of tail-recursion optimizations that are not found in certain strict functional languages⁴ including Standard ML, Scheme, and the concurrent functional language Erlang. However, functions in Oz are not lazy. This property is easily programmed using the concurrency constructs of Oz and the single assignment variables. Here is an example of the well-known higher order function `Map/2`. It is tail recursive in Oz but not in Standard ML or in Scheme.

```

fun {Map Xs F}
  case Xs
  of nil then nil
  [] X|Xr then {F X}|{Map Xr F}
  end
end
{Browse {Map [1 2 3 4] fun {$ X} X*X end}}

```

andthen and orelse

After all, we have been doing a lot of work for nothing! Oz already provides the Boolean lazy (non-strict) versions of the functions `And/2` and `Or/2` as the Boolean operators **andthen** and **orelse** respectively. The former behaves like the function `AndThen/2`, and the latter evaluates its second argument only if the first argument

⁴ Strict functional languages evaluate all its argument first before executing the function.

evaluates **false**. As usual, these operators are not primitives, they are defined in Oz. Figure 11 defines the final version of the function `BinaryTree`.

```
fun {BinaryTree T}
  case T of nil then true
  [] tree(K V T1 T2) then
    {BinaryTree T1} andthen {BinaryTree T2}
  else false end
end
```

- Figure 11 Checking a binary tree lazily.

To Function or not to function? The question is when to use functional notation, and when not. The honest answer is that it is up to you! I will tell you my personal opinion. Here are some rules of thumb:

- First, what I do not like. Given that you defined a procedure `P` do not call it as a function, i.e. do not use functional nesting for procedures. Use instead procedural nesting, with nesting marker, as in the `Merge` example. Moreover, given that you defined a function, call it as function.
- I tend to use function definitions when things are really functional, i.e. there is one output and, possibly many inputs, and the output is a mathematical function of the input arguments.
- I tend to use procedures in most of the other cases, i.e. multiple outputs or nonfunctional definition due to stateful data types or nondeterministic definitions⁵.
- One may relax the previous rule and use functions when there is a clear direction of information-flow although the definition is not strictly functional. Hence, by this rule we can write the binary-merge in Figure 6 as a function although it is definitely not. After all functions are concise.

⁵ In fact, I use mostly objects except for typical constraint based applications.

5 Modules and Interfaces

Modules, also known as packages, are collection of procedures and other values⁶ that are constructed together to provide certain related functionality. A module typically has a number of private procedures that are not visible outside the module and a number of interface procedures that provides the external services of the module. The interface procedures provide the services of the module. In Oz, there is no syntactic support for modules or interfaces. Instead, the lexical scoping of the language and the record data-type suffices to construct modules. The general a module construction looks as follows. Assume that we would like to construct a module called `List` that provides a number of interface procedures for appending, sorting and testing membership of lists. This would look as follows.

```
declare List in
local
  proc {Append ... } ... end
  proc {Sort ... } ... {MergeSort ...} ... end
  proc {Member ...} ... end
  proc {MergeSort ...} ... end
in
  List = list(append: Append
              sort: Sort
              member: Member
              ... )
end
```

Access to `Append` procedure outside of the module `List` is done by using the field `append` from the 'record' `List:List.append`. Notice that in the above example the procedure `MergeSort` is private to the module. Most of the library modules of DFKI Oz follow the above structure.

Often some of the procedures in a module are made global without the module prefix. There are several ways to do this. Here is one way to make `Append` externally visible:

```
declare List Append in
local
  proc {!Append ... } ... end
  ...
in
  List = list(append: Append
              sort: Sort
              member: Member
              ... )
end
```

Modules are often stored on files, for example the `List` module could be stored on the file '`List.oz`'. This file may be inserted into another file by using the compiler directive `\insert 'List.oz'`.

⁶ Classes, objects, etc.

You can investigate the libraries accompanying the system in the EMACS OZ menu. Click on the entry Find, and select Modules File.

The List module is defined in the 'Oz Standard Modules' documentation.

Compiler directives are defined in 'The DFKI Oz User's Manual'.

So far, we have seen only one thread executing. It is time to introduce concurrency. In Oz a new concurrent thread of control is spawned by:

thread *S* **end**

Executing this statement, a thread is forked that runs concurrently with the current thread. The current thread resumes immediately with the next statement. Each nonterminating thread, that is not blocking, will eventually be allocated a time slice of the processor. This means that threads are executed fairly.

However, there are three priority levels: *high*, *medium*, and *low* that determine how often a runnable thread is allocated a time slice. In Oz, a high priority thread cannot starve a low priority one. Priority determines only how large piece of the processor cake a thread can get.

Each thread has a unique name. To get the name of the current thread the procedure `Thread.this/1` is called. Having a reference to a thread, by using its name, enables operations on threads such as terminating a thread, or raising an exception in a thread. Thread operations are defined the standard module `Thread`⁷.

Let us see what we can do with threads. First, remember that each thread is a data flow thread that blocks on data dependency. Consider the following program:

```
declare X0 X1 X2 X3 in
thread
  local Y0 Y1 Y2 Y3 in
    {Browse [Y0 Y1 Y2 Y3]}
    Y0 = X0+1
    Y1 = X1+Y0
    Y2 = X2+Y1
    Y3 = X3+Y2
    {Browse completed}
  end
end
{Browse [X0 X1 X2 X3]}
```

If you input this program and watch the display of the Browser tool, the variables will appear unbound. If you now input the following statements one at a time:

```
X0 = 0
X1 = 1
X2 = 2
X3 = 3
```

you will see how the thread resumes and then suspends again. First when `X0` is bound the thread can execute `Y0 = X0+1` and suspends again because it needs the value of `X1` while executing `Y1 = X1+Y0`, and so on.

⁷ Christian. Schulte, et al, Oz Standard Modules, Oz documentation series.

```

fun {Map Xs F}
  case Xs
  of nil then nil
  [] X|Xr then thread {F X} end | {Map Xr F}
  end
end

```

- Figure 12 A concurrent Map function.

The program shown in Figure 12 defines a concurrent Map function. Notice that ‘thread ... end’ is used here as an expression. Let us discuss the behavior of this program. If we enter the following statements:

```

declare
  F X Y Z
  {Browse thread {Map X F} end}

```

a thread executing Map is created. It will suspend immediately in the case-statement because X is unbound. If we thereafter enter the following statements:

```

X = 1 | 2 | Y
fun {F X} X*X end

```

the main thread will traverse the list creating two threads for the first two arguments of the list, **thread** {F 1} **end**, and **thread** {F 2} **end**, and then it will suspend again on the tail of the list Y. Finally,

```

Y = 3 | Z
Z = nil

```

will complete the computation of the main thread and the newly created thread **thread** {F 3} **end**, resulting in the final list [1 4 9]. The program shown in Figure 13 is a concurrent divide-and-conquer program, which is very inefficient way to compute the ‘Fibonacci’ function. This program creates an exponential number of threads! See how it is easy to create concurrent threads. You may use this program to test how many threads your Oz installation can create. Try

```

{Fib 24}

```

while using the panel program in your Oz menu to see the threads. If it works, try a larger number.

```

fun {Fib X}
  case X
  of 0 then 1
  [] 1 then 1
  else thread {Fib X-1} end + {Fib X-2} end
end

```

- Figure 13 A concurrent Fibonacci function.

** The Oz Panel Showing Thread Creation in {Fib 26 X}

The whole idea of explicit thread creation in Oz is to enable the programmer to structure his/her application in a modular way. Therefore, create threads only when the application need it, and not because concurrency is fun.

6.1 Time

In module Time, we can find a number of useful soft real-time procedures. Among them are:

- {Alarm I ?U} which creates immediately its own thread, and binds U to **unit** after I milliseconds.
- {Delay I} suspends the executing thread for, a least, I milliseconds and then reduces to **skip**.

```

local
  proc {Ping N}
    case N==0 then {Browse 'ping terminated'}
    else {Delay 500} {Browse ping} {Ping N-1} end
  end
  proc {Pong N}
    {For 1 N 1
     proc {$ I} {Delay 600} {Browse pong} end }
    {Browse 'pong terminated'}
  end
in
  {Browse 'game started'}
  thread {Ping 50} end
  thread {Pong 50} end
end

```

- Figure 14 A 'Ping Pong' program.

The program shown in Figure 14 starts two threads, one displays *ping* periodically after 500 milliseconds, and the other *pong* after 600 milliseconds. Some pings will be displayed immediately after each other because of the periodicity difference.

6.2 Stream Communication

The data-flow property of Oz easily enables writing threads that communicate through streams in a producer-consumer pattern. A stream is a list that is created incrementally by one thread (the producer) and subsequently consumed by one or more threads (the consumers). The threads consume the same elements of the stream. For example, the program in Figure 15 is an example of stream communication, where the producer generates a list of numbers and the consumer sums all the numbers.

```

fun {Generator N}
  case N > 0 then N|{Generator N-1}
  else nil end
end
local
  fun {Sum1 L A}
    case L
    of nil then A
    [] X|Xs then {Sum1 Xs A+X}
    end
  end
in fun {Sum L} {Sum1 L 0} end
end

```

- Figure 15 Summing the elements in a list.

Try the program above by running the following program:

```

local L in
  thread L = {Generator 150000} end
  thread {Browse {Sum L}} end
end

```

It should produce the number 11250075000. Let us understand the working of stream communication. A producer incrementally creates a stream (a list) of elements as in the following example where it is producing volvo's. This happens in general in an eager fashion.

```

fun {Producer ...} ... volvo|{Producer ...} ... end

```

The consumer waits on the stream until items arrives, then the items are consumed as in:

```

proc {Consumer Ls ...}
  case Ls of volvo|Lr then 'Consume volvo' ... end
  {Consumer Lr}
end

```

The data flow behavior of the *case-statement* suspends the consumer until the arrival of the next item of the stream. The recursive call allows the consumer to iterate the action over again. The following pattern avoids the use of recursion by using an iterator instead:

```

proc {Consumer Ls ...}
  {ForAll Ls
    proc {$ Item}
      case Item of volvo then
        'Consume volvo' ...
      end
    end}
end

```

Figure 16 shows a simple example using this pattern. The consumer counts the cars received. Each time it receives 1000 cars it prints a message on the display of the Browser.

```

fun {Producer N}
  case N > 0 then
    volvo|{Producer N-1}
  else then nil end
end
proc {Consumer Ls}
  proc {Consumer Ls N}
    case Ls
    of nil then skip
    [] volvo|Lr then
      case N mod 1000 == 0 then
        {Browse 'riding a new volvo'}
      else skip end
      {Consumer Lr N+1}
    else {Consumer Lr N} end
  end
in {Consumer Ls 1}
end

```

- Figure 16 Producing volvo's.

You may run this program using:


```
{Consumer thread {Producer 10000} end}
```

When you feed a statement into the emulator, it is executed in its own thread. Therefore, after feeding the above statement two threads are created. The main one is for the consumer, and the forked thread is for the producer.

Notice that the consumer was written using the *recursive* pattern. Can we write this program using the iterative `ForAll/2` construct? This is not possible because the consumer carries an extra argument `N` that accumulates a result which, is passed to the next recursive call. The argument corresponds to some kind of *state*. In general, there are two solutions. We either introduce a stateful (mutable) data structure, which we will do in Section 7, or define another iterator that passes the state around. In our case, some iterators that fit our needs exist in the module `List`. First, we need an iterator that filters away all items except `volvo`'s. We can use `{Filter Xs P ?Ys}` which outputs in `Ys` all the elements that satisfies the procedure `P/2` used as a Boolean function. The second construct is `{List.forAllInd Xs P}` which is similar to `ForAll`, but `P/2` takes the index of the current element of the list, starting from 1, as its first argument, and the element of the list as its second argument. Here is the program:

```
proc {Consumer Ls}
  fun {IsVolvo X} X == volvo end
  Ls1
in
  thread Ls1 = {Filter Ls IsVolvo} end
  {List.forAllInd Ls1
    proc {$ N X}
      case N mod 1000 == 0 then
        {Browse 'riding a new volvo'}
      else skip end
    end}
  end
```

6.3 Thread Priority and Real Time

Try to run the program using the following statement:

```
{Consumer thread {Producer 5000000} end}
```

Switch on the panel and observe the memory behavior of the program. You will quickly notice that this program does not behave well. The reason has to do with the asynchronous message passing. If the producer sends messages i.e. create new elements in the stream, in a faster rate than the consumer can consume, increasingly more buffering will be needed until the system starts to break down⁸. There are a number of ways to solve this problem. One is to create a bounded buffer between producers and consumers which we will be discussed later. Another way is to control experimentally the rate of thread execution so that the consumers get a larger time-slice than the producers do.

⁸ Ironically in our current research project PERDIO, developing next generation Oz system (an Internet-wide distributed Oz), stream communication across sites works better because of designed flow-control mechanism that suspends producers when the network buffers are full.

The modules `Thread`⁹, and `System`¹⁰, provide a number of operations pertinent to threads. Some of these are summarized in Table 6-1.

| Procedure | Description |
|---|---|
| { <code>Thread.state +T ?A</code> } | Returns current state of <code>T</code> |
| { <code>Thread.suspend +T</code> } | Suspends <code>T</code> |
| { <code>Thread.resume +T</code> } | Resumes <code>T</code> |
| { <code>Thread.terminate +T</code> } | Terminates <code>T</code> |
| { <code>Thread.injectException +T +E</code> } | Raises exception <code>E</code> in <code>T</code> |
| { <code>Thread.this +T</code> } | Returns the current thread <code>T</code> |
| { <code>Thread.setPriority +T +P</code> } | Sets <code>T</code> 's priority |
| { <code>Thread.setThisPriority +P</code> } | Sets current thread's priority |
| { <code>System.get priorities +P</code> } | Gets system-priority ratios |
| { <code>System.set priorities(high:+X medium:+Y)</code> } | Sets system-priority ratios |

• Table 6-1 Thread operations.

DFKI Oz has three priority levels. The system procedure

```
{System.set priorities(high:X medium:Y)}
```

sets the processor-time ratio to $X:1$ between high-priority threads and medium-priority thread. It also sets the processor-time ratio to $Y:1$ between medium-priority threads and low-priority thread. X and Y are integers. So, if we execute

```
{System.set priorities(high:+10 medium:+10)}
```

for each 10 time-slices allocated to runnable high-priority threads, the system will allocate one time-slice for medium-priority threads, and similarly between medium and low priority threads. Within the same priority level, scheduling is fair and round-robin. Now let us make our producer-consumer program work. We give the producer low priority, and the consumer high. We also set the priority ratios to $10:1$ and $10:1$.

```
local L in
  {System.set priorities(high:+10 medium:+10)}
  thread
    {Thread.setThisPriority low}
    L = {Producer 5000000}
  end
  thread
    {Thread.setThisPriority high}
    {Consumer L}
  end
end
```

6.4 Demand-driven Execution

An extreme alternative solution is to make the producer lazy, only producing an item when the consumer requests one. A consumer, in the case, constructs the stream with unbound variables (empty boxes). The producer waits for the unbound variables (empty boxes) to appear on the stream. It then binds the variables (fills the boxes). The general pattern of the producer is as follows.

⁹ Defined in 'Oz standard Modules', Oz documentation series.

¹⁰ Defined in 'Oz User Manual', Oz documentation series.

```

proc {Producer Xs}
  case Xs of X/Xr then
    I in 'Produce I'
    X=I ...
    {Producer Xr}
  end
end

```

The general pattern of the consumer is as follows.

```

proc {Consumer ... Xs}
  X Xr in
    ...
    Xs = X/Xr
    'Consume X'
    ... {Consumer ... Xr}
end

```

The program shown in Figure 17 is a demand driven version of the program in Figure 16. You can run it with very large number of volvo's!

```

local
  proc {Producer L}
    case L of X|Xs then X = volvo {Producer Xs}
    [] nil then {Browse 'end of line'}
    end
  end
  proc {Consumer N L}
    case N==0 then L = nil
    else
      X|Xs = L
    in
      case X of volvo then
        case N mod 1000 == 0 then
          {Browse 'riding a new volvo'}
        else skip end
        {Consumer N-1 Xs}
      else {Consumer N Xs} end
    end
  end
in
  {Consumer 10000000 thread {Producer $} end}
end

```

- Figure 17 Producing volvo's lazily.

Thread Termination-Detection

We have seen how threads are forked using the statement **thread S end**. A natural question that arises is how to join back a forked thread into the original thread of control. In fact, this is a special case of detecting termination of multiple threads, and making another thread wait on that event. The general scheme is quite easy because Oz is a data-flow language.

```

thread  $T_1$   $X_1 = \text{unit}$  end
thread  $T_2$   $X_2 = X_1$  end
...
thread  $T_N$   $X_N = X_{N-1}$  end
{Wait  $X_N$ }
MainThread

```

When All threads terminate the variables $X_1 \dots X_N$ will be merged together labeling a single box that contains the value **unit**. {Wait X_N } suspends the main thread until X_N is bound.

In Figure 18 we define a higher-order construct (combinator), that implements the concurrent-composition control construct that has been outlined above. It takes a single argument that is a list of nullary procedures. When it is executed, the procedures are forked concurrently. The next statement is executed only when all procedures in the list terminate.

```

local
  proc {Concl Ps I O}
    case Ps of P|Pr then
      M in
        thread {P} M = I end
        {Concl Pr M O}
      [] nil then O = I
    end
  end
in
  proc {Conc Ps} {Wait {Concl Ps unit $}} end
end

```

- Figure 18 Concurrent Composition.

7 Stateful Data Types

Oz provides set of stateful data types. These include ports, objects, arrays, and dictionaries (hash tables). These data types are abstract in the sense that they are characterized only by the set of operations performed on the members of the type. Their implementation is always hidden, and in fact different implementations exist but their corresponding behavior remains the same. For example, objects are implemented in a totally different way depending on the optimization level of the compiler. Each member is always unique by conceptually tagging it with an Oz-name upon creation. A member is created by an explicit creation operation. A type test operation always exists. In addition, a member ceases to exist when it is no longer accessible.

7.1 Ports

Port is such an abstract data-type. A Port P is an asynchronous communication channel that can be shared among several senders. A port has a stream associated with it. The operation: $\{\text{NewPort } S \ ?P\}$ creates a port P and initially connects it to the variable S taking the role of a stream. The operation: $\{\text{Send } P \ M\}$ will append the message M to the end of the stream associated with P . The port keeps track of the end of the stream as its next insertion point. The operation $\{\text{IsPort } P \ ?B\}$ checks whether P is a port. The following program shows a simple example using ports: X

```
declare S P
P = {Port.new S}
{Browse S}
{Send P 1}
{Send P 2}
```

If you enter the above statements incrementally you will observe that S gets incrementally more defined.

```
S
1 |
1 | 2 | _
```

Ports are more expressive abstractions than pure stream communication that was discussed in Section 6.2, since they can be shared among multiple thread, and can be embedded in other data structures. Ports are the main message passing mechanism between threads in Oz.

Server-Clients Communication

The program shown in Figure 19 defines a thread that acts as FIFO queue server. Using single-assignment (logic) variables makes the server insensitive to the arrival order of `get` messages relative `put` messages. `get` messages can arrive even when the queue is empty. A server is created by $\{\text{NewQueueServer } ?Q\}$. This procedure returns back a port Q to the server. A client thread having access to Q can request services by sending a message on the port. Notice how results are returned back through logic variables. A client requesting an Item in the queue will send the message $\{\text{Send } Q \ \text{get}(I)\}$. The server will eventually answer back by binding I to an item.

```

declare NewQueueServer in
local
  fun {NewQueue}
    X in q(front:X rear:X number:0)
  end
  fun {Put Q I}
    X in
      Q.rear = I|X
      {AdjoinList Q [rear#X number#(Q.number+1)]}
  end
  proc {Get Q ?I ?NQ}
    X in
      Q.front = I|X
      NQ = {AdjoinList Q [front#X number#(Q.number-1)]}
  end
  proc {Empty Q ?B ?NQ}
    B = (Q.number == 0) NQ = Q
  end
  proc {QServer Xs Q}
    case Xs of X|Xr then
      NQ in
        case X
          of put(I) then NQ = {Put Q I}
          [] get(I) then {Get Q I NQ}
          [] empty(B) then {Empty Q B NQ}
          else NQ = Q end
          {QServer Xr NQ}
        [] nil then skip
      end
    end
    S P = {NewPort S}
in fun {NewQueueServer}
  thread {QServer S {NewQueue}} end
  P
end
end

```

- Figure 19 Concurrent Queue server.

The following sequence of statement illustrates the working of the program.

```

declare
Q = {NewQueueServer}
{Send Q put(1)}
{Browse {Send Q get($)}}
{Browse {Send Q get($)}}
{Browse {Send Q get($)}}
{Send Q put(2)}
{Send Q put(3)}
{Browse {Send Q empty($)}}

```

*** Explain the use of logic variables as a mechanism to returns value.

*** Explain the common pattern of using nesting markers in message passing.

7.2 Chunks

Ports are actually stateful data structures. A port keeps a local state internally tracking the end of its associated stream. Oz provides two primitive devices to construct abstract stateful data-types *chunks* and *cells*. All others subtypes of chunks can be defined in terms of chunks and cells.

A chunk is similar to a record except that the label of a chunk is an oz-name, and there is no arity operation available on chunks. This means one can hide certain components of a chunk if the feature of the component is an oz-name that is visible only (by lexical scoping) to user-define operations on the chunk.

A chunk is created by the procedure `{NewChunk Record}`. This creates a chunk with the same arguments as the record, but having a unique label. The following program creates a chunk.

```
local X in
  {Browse X={NewChunk f(c:3 a:1 b:2)}}
  {Browse X.c}
end
```

This will display the following.

```
<Ch>(a:1 b:2 c:3)
3
```

In Figure 20, we show an example of using the information hiding ability of chunks to implement Ports.

7.3 Cells

A cell could be seen as a chunk with a mutable single component. A cell is created as follows.

```
{NewCell X ?C}
```

A cell is created with the initial content `x`. `C` is bound to a cell. The `Table{Cell}` shows the operations on a cell.

| Operation | Description |
|-------------------|---|
| {NewCell X ?C} | Creates a cell C with content X. |
| {Access +C X} | Returns the content of C in X. |
| {Assign +C Y} | Modifies the content of C to Y. |
| {IsCell +C} | Tests if C is a cell |
| {Exchange +C X Y} | Swaps atomically the content of C from X to Y |

- Table 7-1 Cell operations.

Check the following program. The last statement increments the cell by one. If we leave out **thread** ... **end** the program deadlocks. Do you know why?

```

local I O X in
  I = {NewCell a} {Browse {Access I}}
  {Assign I b}      {Browse {Access I}}
  {Assign I X}      {Browse {Access I}}
  X = 5*5
  {Exchange I O thread O+1 end} {Browse {Access I}}
end

```

Cells and higher-order iterators allow conventional assignment-based programming in

Oz. The following program accumulates in the cell J $\sum_{i=1}^{10} i$.

```

declare J in
  J = {NewCell 0}
  {For 1 10 1
    proc {$ I}
      O N in
        {Exchange J O N}
        N = O+I
      end}
    {Browse {Access J}}
  }

```

Ports described in Subsection 7.1 can be implemented by chunks and cells in a secure way, i.e. as an abstract data type that cannot be forged. The following program shows an implementation of Ports.

```

declare NewPort IsPort Send in
local
  Port = {NewName} %New Oz name
in
  fun {NewPort S}
    C = {NewCell S}
    {NewChunk port(Port:C)}
  end
  fun {IsPort ?P}
    {ChunkHasFeature Port} %Checks a chunk feature
  end
  proc {Send P M}
    Ms Mr in
      {Exchange P.Port Ms Mr}
      Ms = M|Mr
  end

```

```
end
end
```

- Figure 20. Implementation of Ports by Cells and Chunks.

Initially an Oz-name is created locally, which is accessible only by the Port operations. A port is created as a chunk that has one component, which is a cell. The cell is initialized to the stream associated with the port. The type test `IsPort` is done by checking the feature `Port`. Sending a message to a port results in updating the stream atomically, and updating the cell to point to the tail of the stream.

8 Classes and Objects

A Class in Oz is a chunk that contains:

- A collection of methods in a method table.
- A description of the attributes that each instance of the class will possess. Each attribute is a stateful cell that is accessed by the attribute-name, which is either an atom or an Oz-name.
- A description of the features that each instance of the class will possess. A feature is an immutable component (a variable) that is accessed by the feature-name, which is either an atom or an Oz-name.
- Classes are stateless Oz-values¹¹. Contrary to languages like Smalltalk, or Java etc., they are just descriptions of how the objects of the class should behave.

Classes from First Principles

Figure 21 shows how a class is constructed from first principles as outlined above. Here we construct a Counter class. It has a single attribute accessed by the atom `val`. It has a method table, which has three methods accessed through the chunk features `browse`, `init` and `inc`. A method is a procedure that takes a message, always a record, an extra parameter representing the state of the current object, and the object itself known internally as `self`.

```
declare Counter
local
  Attrs = [val]
  MethodTable = m(browse:MyBrowse init:Init inc:Inc)
  proc {Init M S Self}
    init(Value) = M in
      {Assign S.val Value}
    end
  end
  proc {Inc M S Self}
    X inc(Value)=M in
      {Access S.val X} {Assign S.val X+Value}
    end
  end
  proc {MyBrowse M=browse S Self}
    {Browse {Access S.val}}
  end
in
  Counter = {NewChunk c(methods:MethodTable attrs:Attrs)}
end
```

- Figure 21. An Example of Class construction.

*** Talk about the example.

Objects from First Principles

Figure 22 shows a generic procedure that creates an object from a given class. This procedure creates an object state from the attributes of the class. It initializes the attributes of the object, each to a cell (with unbound initial value). We use here the

¹¹ In fact, classes may have some invisible state. In the current implementation, a class usually has method cache, which is stateful.

iterator `Record.forAll/2` that iterates over all fields of a record. `NewObject` returns a procedure `Object` that identifies the object. Notice that the state of the object is visible only within `Object`. One may say that `Object` is a procedure that encapsulates the state¹².

```

proc {NewObject Class InitialMethod ?Object}
  local
    State O
  in
    State = {MakeRecord s Class.attrs}
    {Record.forAll State proc {$ A} {NewCell A} end}
    proc {O M}
      {Class.methods.{Label M} M State O}
    end
    {O InitialMethod}
    Object = O
  end
end

```

- Figure 22. Object Construction.

We can try our program as follows

```

declare C
  {NewObject Counter init(0) C}
  {C inc(6)} {C inc(6)}
  {C browse}

```

Try to execute the following statement.

```

local X in {C inc(X)} X=5 end {C browse}

```

You will see that nothing happens. The reason is that the object application

```

{C inc(X)}

```

suspends inside the procedure `INC/3` that implements method `inc`. Do you know where exactly? If you on the other hand execute the following statement, things will work as expected.

```

local X in thread {C inc(X)} end X=5 end {C browse}

```

8.1 Objects and Classes for Real

Oz supports object-oriented programming following the methodology outlined above. There is also syntactic support and optimized implementation so that object application (calling a method in objects) is as cheap as procedure calls. The class `Counter` defined earlier has the syntactic form shown in Figure 23:

```

class Counter
  attr val
  meth browse
    {Browse @val}
  end
  meth inc(Value)
    val <- @val + Value
  end

```

¹² This is a simplification; an object in Oz is a chunk that has the above procedure in one of its fields; other fields contain the object features.

```

    meth init(Value)
        val <- Value
    end
end

```

- Figure 23. Counter Class.

A class X is defined by:

```
class X ... end.
```

Attributes are defined using the attribute-declaration part before the method-declaration part:

```
attr  $A_1 \dots A_N$ 
```

Then follows the method declarations, each has the form:

```
meth  $E \ S \ end$ 
```

where the expression E evaluates to a method head, which is a record whose label is the method name. An attribute A is accessed using the expression $@A$. It is assigned a value using the statement $A <- E$.

A class can be defined anonymously by:

```
 $X = class \ \$ \dots end.$ 
```

The following shows how an object is created from a class using the procedure `New/3`, whose first argument is the class, the second is the initial method, and the result is the object. `New/3` is a generic procedure for creating objects from classes.

```

declare C in
C = {New Counter init(0)}
{C browse}
{C inc(1)}
local X in thread {C inc(X)} end X=5 end

```

8.2 Static Method Calls

Given a class C and a method head $m(\dots)$, a method call has the following form:

$C \ , \ m(\dots)$

A method call invokes the method defined in the class argument. A method call can only be used inside method definitions. This is because a method call takes the current object denoted by `self` as implicit argument. The method could be defined the class or inherited from a super class. Inheritance will be explained shortly.

Classes as Modules

Static method calls have in general the same efficiency as procedure calls. This allows classes to be used as modules. This is advantageous because classes can be built incrementally by inheritance. The program shown in Figure shows a possible class acting as a module. The class `ListC` defines some common list-procedures as methods. `ListC` defines the methods `append/3`, `member/2`, `length/2`, and `nrev/2`. Notice that a method body is similar to any Oz statement but in addition, method calls are allowed. We also see the first example of inheritance.

Inheritance is a way to construct new classes from existing classes. It defines what attributes, features¹³, and methods are available in the new class. We will restrict our discussion of inheritance to methods. Nonetheless, the same rules apply to features and attributes.

The methods are available in a class *C* (i.e. visible) are defined through a precedence relation on the methods that appear in the class hierarchy. We call this relation the *overriding relation*:

- A method in a class *C* overrides any method, with the same label, in any super class if *C*.
- A method defined in a class *B* declared in the **from**-declaration of *C* overrides any method, with the same label, defined in a class to the left of *B* in the **from** declaration.

Now a class hierarchy with the super-class relation can be seen as a directed graph with the class being defined as the root. The edges are directed towards the subclasses. There are two requirements for the inheritance to be valid. First, the inheritance relation is directed and acyclic. So the following is not allowed:


```
class A from B ... end
class B from A ... end
```



• Figure 25. Illegal class hierarchy.

Second, after striking out all overridden methods each remaining method should have a unique label and is defined only in one class in the hierarchy. Hence, class *C* in the following example is not valid because the two methods labeled *m* remains.

```
class A meth m(...) ... end end
class B meth m(...) ... end end
class B from B1 end
class A from A1 end
class C from A B end
```



• Figure 26. Illegal class hierarchy in method *m*.

Whereas the class *C* below is valid, and the method *m* that is available in *C* is that of *B*.

¹³ To be defined shortly.


```

class A meth m(...) ... end end
class B meth m(...) ... end end
class C from A B end

```

Notice that if you run a program with an invalid hierarchy, the system will not complain until an object is created that tries to access an invalid method. Only at this point of time, you are going to get a runtime exception. The reason is that classes are partially formed at compile time, and are completed by demand, using method caches, at execution time.

Multiple inheritance or Not

My opinion is the following:

- In general, to use multiple inheritance correctly, one has to understand the total inheritance hierarchy, which is sometimes worth the effort.
- I do not like the existence of any overriding rule that depends on the order of classes in the **from** construct. So I would consider the latter example as invalid as the former one. The reason is that if you take **A** and **B** in the later example, and refine them to get the first example your working program will suddenly cease to work.
- If there is a method-name conflict between immediate super classes, I would define the method locally to overrides the conflict-causing methods.
- There is another problem with multiple inheritance when sibling super-classes share (directly or indirectly) a common ancestor-class that is stateful (i.e. has attributes). One may get replicated operations on the same attribute. This typically happens when executing an initialization method in a class, one has to initialize its super classes. The only remedy here is to understand carefully the inheritance hierarchy to avoid such replication. Alternatively, you should only inherit from multiple classes that do not share stateful common ancestor. This problem is known as the implementation-sharing problem.

8.4 Features

Objects may have features similar to records. Features are components that are specified in the class declaration:

```

class C from ...
  feat  $a_1 \cdots a_N$ 
...
end

```

As in a record, a feature of an object has an associated field. The field is a logic variable that can be bound to any Oz value (including cells, objects, classes etc.). Features of objects are accessed using the infix ``.´` operator. The following shows an example using features:

```

class ApartmentC from BaseObject
  meth init skip end
end
class AptC from ApartmentC
  feat
    streetName: york
    streetNumber:100
    wallColor:white
    floorSurface:wood
  end
end

```

Feature initialization

The example shows how features could be initialized at the time the class is defined. In this case, all instances of the class AptC will have the features of the class, with their corresponding values. Therefore, the following program will display *york* twice.

```

declare Apt1 Apt2
Apt1 = {New AptC init}
Apt2 = {New AptC init}
{Browse Apt1.streetName}
{Browse Apt2.streetName}

```

We may leave a feature uninitialized as in:

```

class MyAptC1 from ApartmentC
  feat streetName
end

```

In this case whenever an instance is created, the field of the feature is assigned a new fresh variable. Therefore, the following program will bind the feature *streetName* of object Apt3 to the atom *kungsgatan*, and the corresponding feature of Apt4 to the atom *sturegatan*.

```

declare Apt3 Apt4
Apt3 = {New MyAptC1 init}
Apt4 = {New MyAptC1 init}
Apt3.streetName = kungsgatan
Apt4.streetName = sturegatan

```

One more form of initialization is available. A feature may be initialized in the class declaration to a variable or an Oz-value that has a variable. In the following, the feature is initialized to a tuple with an anonymous variable. In this case, all instances of the class will *share* the same variable. Consider the following program.

```

class MyAptC1 from ApartmentC
  feat streetName:f(____)
end

local Apt1 Apt2 in
  Apt1 = {New MyAptC1 init}
  Apt2 = {New MyAptC1 init}
  {Browse Apt1.streetName}
  {Browse Apt2.streetName}
  Apt1.streetName = f(york)
end

```

If entered incrementally, will show that the statement

```
Apt1.streetName = f(york)
```

binds the corresponding feature of `Apt2` to the same value as that of `Apt1`.

8.5 Parameterized Classes

There are many ways to get your classes more generic, which later may be specialized for specific purposes. The common way to do this in object-oriented programming is to define first *an abstract class* in which some methods are left unspecified. Later these methods are defined in the subclasses. Suppose you have defined a generic class for sorting where the comparison operator `less` is needed. This operator depends on what kinds of data are being sorted. Different realizations are needed for integer, rational, or complex numbers, etc. In this case, by subclassing we can specialize the abstract class to a 'concrete' class.

In Oz, we have also another natural method for creating generic classes. Since classes are first-class values, we can instead define a function that takes some type argument(s) and return a class that is specialized for the type(s). In Figure 27, the function `SortClass` is defined that takes a class as its single argument and returns a sorting class specialized for the argument.

```
fun {SortClass Type}
  class $ from BaseObject
    meth qsort(Xs Ys)
      case Xs
      of nil then Ys = nil
      [] P|Xr then S L in
        {self partition(Xr P S L)}
        ListC, append(C,qsort(S $) P|(C,qsort(L $)) Ys)
      end
    end
    meth partition(Xs P Ss Ls)
      case Xs
      of nil then Ss = nil Ls = nil
      [] X|Xr then Sr Lr in
        case Type,less(X P $) then
          Ss = X|Sr Lr = Ls
        else
          Ss = Sr Ls = X|Lr
        end
        C,partition(Xr P Sr Lr)
      end
    end
  end
end
```

• Figure 27. Parameterized Classes.

We can now define two classes for integers and rationals:

```

class Int
  meth less(X Y B)
    B = X<Y
  end
end
class Rat from Object
  meth less(X Y B)
    '/'(P Q) = X
    '/'(R S) = Y
  in
    B = P*S < Q*R
  end
end

```

Thereafter, we can execute the following statements:

```

{Browse {{New {SortClass Int} noop} qsort([1 2 5 3 4] $)}}
{Browse {{New {SortClass Rat} noop}
  qsort(['/'(23 3) '/'(34 11) '/'(47 17)] $)}}

```

8.6 Self Application

The program in Figure 27 shows in the method `qsort` an object application using the keyword **self** (see below).

```

meth qsort(Xs Ys)
  case Xs
  ...
  {self partition(Xr P S L)}
  ...
end

```

We use here the phrase *object-application* instead of the commonly known phrase *message sending* because message sending is misleading in a concurrent language like Oz. When we use **self** instead of a specific object as in

```
{self partition(Xr P S L)}
```

we mean that we dynamically pick the method `partition` that is defined (available) in the current object. Thereafter we apply the object (as a procedure) on the message. This is a form of dynamic binding common in all object-oriented languages.

8.7 Attributes

We have touched before on the notion of attributes. Attributes are the carriers of state in objects. Attributes are declared similar to features, but using the keyword **attr** instead. When an object is created each attribute is assigned a new cell as its value. These cells are initialized very much the same way as features. The difference lies in the fact that attributes are cells that can be assigned, reassigned and accessed at will. However, attributes are private to their objects. The only way to manipulate an attribute from outside an object is to force the class designer to write a method that manipulates the attribute. In the Figure{Point} we define an the class `Point`. Note that the attributes `x` and `y` are initialized to zero before the initial message is applied. The method `move` uses **self**-application internally.

```

class Point from BaseObject
  attr x:0 y:0
  meth init(X Y)
    x <- X
    y <- Y          % attribute update
  end
  meth location(L)
    L = l(x:@x y:@y) % attribute access
  end
  meth moveHorizontal(X)
    x <- X
  end
  meth moveVertical(Y)
    y <- Y
  end
  meth move(X Y)
    {self moveHorizontal(X)}
    {self moveVertical(Y)}
  end
  meth display
    % Switch the browser to virtual string mode
    {Browse "point at ("#@x#" , "#@y#")\n"}
  end
end

```

- Figure 28. The class Point.

Try to create an instance of Point and apply some few messages:

```

declare P
P = {New Point init(2 0)}
{P display}
{P move(3 2)}

```

8.8 Private and Protected Methods

Methods may be labeled by variables instead of literals. These methods are *private* to the class in which they are defined, as in:

```

class C from ...
  meth A(X) ... end
  meth a(...) {self A(5)} ... end
  ....
end

```

The method A is visible only within the class C. In fact the notation above is just an abbreviation of the following expanded definition:

```

local A = {NewName} in
  class C from ...
    meth A(X) ... end
    meth a(...) {self A(5)} ... end
    ...
  end
end

```

where A is bound to a new name in the lexical scope of the class definition.

Some object-oriented languages have also the notion of protected methods. A method is *protected* if it is accessible only in the class it is defined or in descendant classes, i.e. subclasses and subclasses etc. In Oz there is no direct way to define a method to be protected. However there is a programming technique that gives the same effect. We know that attributes are only visible inside a class or to descendants of a class by inheritance. We may make a method protected by firstly making it private and secondly storing it in an attribute. Consider the following example:

```
class C from ...
  attr pa:A
  meth A(X) ... end
  meth a(...) {self A(5)} ... end
  ...
end
```

Now, we create a subclass C1 of C and access method A as follows:

```
class C1 from C
  meth b(...) {self @pa(5)} ... end
  ...
end
```

Method b accesses method A through the attribute pa.

Let us continue our simple example in Figure 28 by defining a specialization of the class that in addition of being a point, it stores a history of the previous movement. This is shown in Figure 29.

```
class HistoryPoint from Point
  attr
    history: nil
    displayHistory: DisplayHistory
  meth init(X Y)
    Point,init(X Y) % call your super
    history <- [l(X Y)]
  end
  meth move(X Y)
    Point,move(X Y)
    history <- l(X Y)|@history
  end
  meth display
    Point,display
    {self DisplayHistory}
  end
  meth DisplayHistory % made protected method
    {Browse "with location history: "}
    {Browse @history}
  end
end
```

• Figure 29. The class History Point.

There are a number of remarks on the class definition HistoryPoint. First observe the typical pattern of method refinement. The method move specializes that of class Point. It first calls the super method, and then does what is specific to being a HistoryPoint class. Second, DisplayHistory method is made private to the class. Moreover it is made available for subclasses, i.e. protected, by storing it in the attribute displayHistory. You can now try the class by the following statements:

```
declare P
P = {New HistoryPoint init(2 0)}
{P display}
{P move(3 2)}
```

8.9 Default Argument Values

A method head may have default argument values. Consider the following example.

```
meth m(X Y d1:Z<=0 d2:W<=0) ... end
```

A call of the method `m` may leave the arguments of features `d1` and `d2` unspecified. In this case these arguments will assume the value zero.

We continue our `Point` example by specializing `Point` in a different direction. We define the class `BoundedPoint` as a point that moves in a constrained rectangular area. Any attempt to move such a point outside the area will be ignored. The class is shown in `{BoundedPoint}`. Notice that the method `init` has two default arguments that give a default area if not specified in the initialization of a new instance of `BoundedPoint`.

```

class BoundedPoint from Point
  attr
    xbounds: 0#0
    ybounds: 0#0
    boundConstraint: BoundConstraint
  meth init(X Y xbounds:XB <= 0#10 ybounds:YB <= 0#10)
    Point,init(X Y) % call your super
    xbounds <- XB
    ybounds <- YB
  end
  meth move(X Y)
    case {self BoundConstraint(X Y $)} then
      Point,move(X Y)
    else skip end
  end
  meth BoundConstraint(X Y B)
    B = (X >= @xbounds.1 andthen
        X <= @xbounds.2 andthen
        Y >= @ybounds.1 andthen
        Y <= @ybounds.2 )
  end
  meth display
    Point,display
    {self DisplayBounds}
  end
  meth DisplayBounds
    X0#X1 = @xbounds
    Y0#Y1 = @ybounds
    S = "xbounds=( "#X0#", "#X1#" ), ybounds=( "
        #Y0#", "#Y1#" )"
  in
    {Browse S}
  end
end

```

• Figure 30 The class BoundedPoint

We conclude this section by finishing our example in a way that shows the multiple inheritance problem. We would like now a specialization of both `HistoryPoint` and `BoundedPoint` as a bounded-history point. A point that keeps track of the history and moves in a constrained area. We do this by defining the class `BHPoint` that inherits from the two previously defined classes. Since they both share the class `Point`, which contains stateful attributes, we encounter the implementation-sharing problem. We, any way, anticipated this problem and therefore created two protected methods stored in `boundConstraint` and `displayHistory` to avoid repeating the same actions. In any case, we have to refine the methods `init`, `move`, and `display` since they occur in the two sibling classes. The solution is shown in `{BHPoint}`. Notice how we use the protected methods. We did not care avoiding the repetition of initializing the attributes `x?` and `y?` since it does not make so much harm. Try the following example:

```

declare P
P = {New BHPoint init(2 0)}
{P display}
{P move(1 2)}

```

This pretty much covers most of object system. What is left is how to deal with concurrent threads sharing common space of objects.


```

class BHPoint from HistoryPoint BoundedPoint
  meth init(X Y xbounds:XB <= 0#10 ybounds:YB <= 0#10)
    % repeats init
    HistoryPoint,init(X Y)
    BoundedPoint,init(X Y xbounds:XB ybounds:YB)
  end
  meth move(X Y)
    L = @boundConstraint in
    case {self L(X Y $)} then
      HistoryPoint,move(X Y)
    else skip end
  end
  meth display
    BoundedPoint,display
    {self @displayHistory}
  end
end

```

- Figure 31 The class BHPoint.

9 Objects and Concurrency

As we have seen, objects in Oz are stateful data structures. Threads are the active computation entities. Threads can communicate either by message passing using ports, or through common shared objects. Communication through shared objects requires the ability to serialize concurrent operations on objects so that the object state is kept coherent after each such an operation. In Oz, we separate the issue of acquiring exclusive access of an object from the object system. This gives us the ability to perform coarse-grain atomic operation of a set of objects, a very important requirement in distributed database system. The basic mechanism in Oz to get exclusive access is through locks.

9.1 Locks

The purpose of a lock is to mediate exclusive access to a shared resource between threads. Such a mechanism is typically made safer and more robust by restricting this exclusive access to a critical region. On entry into the region, the lock is secured and the thread is granted exclusive access rights to the resource, and when execution leaves the region, whether normally or through an exception, the lock is released. A concurrent attempt to obtain the same lock will block until the thread currently holding it has released it.

9.1.1 Simple Locks

In the case of a simple lock, a nested attempt by the same thread to reacquire the same lock during the dynamic scope of a critical section guarded by the lock will block. We say *reentrancy* is not supported. Simple locks can be modeled in Oz as follows, where `Code` is a nullary procedure encapsulating the computation to be performed in the critical section. The lock is represented as a procedure, which when applied to some code it tries to get the lock by waiting until `Old` gets bound to `unit`. Notice that the lock is released upon normal as well as abnormal exit.

```
proc {NewSimpleLock ?Lock}
  Cell = {NewCell unit}
in
  proc {Lock Code}
    Old New in
      try
        {Exchange Cell Old New}
        {Wait Old} {Code}
      finally New=unit end
    end
  end
end
```

Atomic Exchange on Object Attributes

Another implementation is using an object as shown below to implement a lock. Notice the use of the construct:

```
Old = lck <- New
```

Similar to the `Exchange` operation on cells, this is an atomic exchange on an object attribute.

```

class SimpleLock
  attr lck:unit
  meth init skip end
  meth lock(Code)
    Old New in
      try
        Old = lck <- New
        {Wait Old} {Code}
      finally New= unit end
    end
  end
end

```

9.1.2 Thread-Reentrant Locks

In Oz, the computational unit is the thread. Therefore an appropriate locking mechanism should grant exclusive access rights to threads. As a consequence the non-reentrant simple lock mechanism presented above is inadequate. A thread-reentrant lock allows the same thread to reenter the lock, i.e. to enter a dynamically nested critical region guarded by the same lock. Such a lock can be acquired by at most one thread at a time. Concurrent threads that attempt to get the same lock are queued. When the lock is released, it is granted to the thread standing first in line etc. Thread-reentrant locks can be modeled in Oz as follows:

```

class ReentrantLock from SimpleLock
  attr Current:unit
  meth lck(Code)
    ThisThread = {Thread.this} in
      case ThisThread == @Current then
        {Code}
      else
        try
          Code1 = proc {$}
            Current <- ThisThread
            {Code}
          end
          in SimpleLock, lck{Code1}
        finally Current <- unit end
      end
    end
  end
end

```

Thread reentrant locks are given syntactic and implementational support in Oz. They are implemented as subtype chunks. Oz provides the following syntax for guarded critical regions:

```
lock E then S end
```

where *E* is an expression that evaluates to a lock. The construct blocks until *S* is executed. If *E* is not a lock, then a type error is raised.

- {NewLock *L*} creates a new lock *L*.
- {IsLock *E*} returns true iff *E* is a lock..

Arrays

Oz has arrays as chunk subtype. Operations on arrays are defined in module Array.

- `{NewArray +L +H +I ?A}` creates an array A, where L is the lower-bound index, H is the higher-bound index, and I is the initial value of the array elements.
- `{Array.low +A ?L}` returns the lower index.
- `{Array.high +A ?L}` returns the higher index.
- `{Get +A +I ?R}` returns A[I] in R.
- `{Put +A +I X}` assigns X to the entry A[I].

As a simple illustration of the use of locks consider the program in Figure 32. The procedure `Switch` transforms negative elements of an array to positive, and zero elements to the atom 'zero'. The procedure `Zero` resets all elements to zero.

```

declare A L in
A = {NewArray 1 100 ~5}
L = {NewLock}
proc {Switch A}
  {For {Array.low A} {Array.high A} 1
    proc {$ I}
      X = {Get A I} in
        case X<0 then {Put A I ~X}
        elsecase X == 0 then {Put A I zero} else skip end
        {Delay 100}
      end}
  end
proc {Zero A}
  {For {Array.low A} {Array.high A} 1
    proc {$ I} {Put A I 0} {Delay 100} end}
  end

```

- Figure 32 Using Lock.

Try the following program.

```

local X Y in
  thread {Zero A} X = unit end
  thread {Switch A} Y = X end
  {Wait Y}
  {For 1 10 1 proc {$ I} {Browse {Get A I}} end}
end

```

The elements of the array will be mixed 0 and 'zero'.

Assume that we want to perform the procedures `Zero` and `Switch`, each atomically but in an arbitrary order. To do this we can use locks as in the following example.

```

local X Y in
  thread
    {Delay 100}
    lock L then {Zero A} end
    X = unit
  end
  thread
    lock L then {Switch A} end
    Y = X
  end
  {Wait Y}
  {For 1 10 1 proc {$ I} {Browse {Get A I}} end}
end

```

By Switching the delay statement above between the first and the second thread, we observe that all the elements of the array either will get the value zero or 0. We have no mixed values.

*** Write an example of an atomic transaction on multiple objects using multiple locks.

9.2 Locking Objects

To guarantee mutual exclusion on objects one may use the locks described in the previous subsection. Alternatively, we may declare in the class that its instance objects can be locked with a default lock existing in the objects when they are created. A class with an implicit lock is declared as follows:

```

class C from ....
  prop locking
  ....
end

```

This does not automatically lock the object when one of its methods is called. Instead we have to use the construct:

```
lock S end
```

inside any method to guarantee exclusive access when *S* is executed. Remember that our locks are thread-reentrant. This implies that:

- if we take all objects that we have constructed and enclose each method body with **lock... end**, and
- execute our program with only one thread, then
- the program will behave exactly as before

Of course, if we use multiple threads calling methods in multiple objects, we might deadlock if there is any cyclic dependency. Writing nontrivial concurrent program needs careful understanding of the dependency patterns between threads. In such programs deadlock may occur whether locks are used or not. It suffices to have a cyclic communication pattern for deadlock to occur.

The program in Figure 23 can be refined to work in concurrent environment by refining it as follows:

```

class CCounter from Counter
  prop locking
  meth inc(Value)
    lock Counter, inc(Value) end
  end
  meth init(Value)
    lock Counter, init(Value) end
  end
end

```

Let us now study a number of interesting examples where threads not only perform atomic transactions on objects, but also synchronize through objects.

Concurrent FIFO Channel

The first example shows a concurrent channel, which is shared among an arbitrary number of threads. Any producing thread may put information in the channel asynchronously. A consuming thread has to wait until information exists in the channel. Waiting threads are served fairly. Figure 33 shows one possible realization. This program relies on the use of logical variables to achieve the desired synchronization. The method `put/1` inserts an element in the channel. A thread executing the method `get/1` will wait until an element is put in the channel. Multiple consuming threads will reserve their place in the channel, thereby achieving fairness. Notice that `{wait I}` is done outside an exclusive region. If waiting was done inside `lock ... end` the program would deadlock. So, as a rule of thumb:

- Do not wait inside an exclusive region, if the waking-up action has to acquire the same lock..

```

class Channel from BaseObject
  prop locking
  attr f r
  meth init
    X in f <- X r <- X
  end
  meth put(I)
    X in lock @r=I | X r<-X end
  end
  meth get(?I)
    X in lock @f=I | X f<-X end {Wait I}
  end
end

```

• Figure 33 An Asynchronous Channel Class

Monitors

The next example shows a traditional way to write *monitors*. We start by defining a class that defines the notion of events and the monitor operations `notify(Event)` and `wait(Event)` by specializing the class `Channel`.

```

class Event from Channel
  meth wait
    Channel , get(_)
  end
  meth notify
    Channel , put(unit)
  end
end

```

We show here an example of a unit buffer in the traditional monitor style. The unit buffer behaves in a way very similar to a channel when it comes to consumers. Each consumer waits until the buffer is full. In the case of producers only one is allowed to insert an item in the empty buffer. Other producers have to suspend until the item is consumed. The program in Figure 34 shows a single buffer monitor. Here we had to program a signaling mechanism for producers and consumers. Observe the pattern in `put/1` and `get/1` methods. Most execution is done in an exclusive region. If waiting is necessary it is done outside the exclusive region. This is done by using an auxiliary variable `x`, which gets bound to `yes`. The `get/1` method notifies one producer at a time by setting the `empty` flag and notifying one producer (if any). This is done as an atomic step. The `put/1` method does the reciprocal action.

```

class UnitBufferM
  attr item empty psignal csignal
  prop locking
  meth init
    empty <- true
    psignal <- {New Event init}
    csignal <- {New Event init}
  end
  meth put(I)
    X in
    lock
      case @empty then
        item <- I
        empty <- false
        X = yes
        {@csignal notify}
      else X = no end
    end
    case X == no then
      {@psignal wait}
      {self put(I)}
    else skip end
  end
  meth get(I)
    X in
    lock
      case {Not @empty} then
        I = @item
        empty <- true
        {@psignal notify}
        X = yes
      else X = no end
    end
    case X == no then
      {@csignal wait}
      {self get(I)}
    else skip end
  end
end

```

• Figure 34 A Unit Buffer Monitor

Try the above example by running the following code:

```

local
  UB = {New UnitBufferM init} in
  {For 1 15 1
    proc{$ I} thread {UB put(I)} {Delay 500} end end}
  {For 1 15 1
    proc{$ I} thread {UB get({Browse})}{Delay 500} end end}
end

```

Bounded Buffers Oz Style

In Oz, it is very rare to write programs in the monitor style shown above. In general it is very awkward. There is a simpler way to write a UnitBuffer class that is not traditional. This is due to the combination of objects and logic variable Figure 35 shows simple definition. No locking is needed directly.


```

class UnitBuffer from BaseObject
  attr prodq buffer
  meth init
    buffer <- {New Channel init}
    prodq <- {New Event init}
    {@prodq notify}
  end
  meth put(I)
    {@prodq wait}
    {@buffer put(I)}
  end
  meth get(?I)
    {@buffer get(I)}
    {@prodq notify}
  end
end

```

- Figure 35 Unit Buffer.

Simple generalization of the above program leads to an arbitrary size bounded buffer class. This is shown in below. The put and get methods are the same as before. Only the initialization method is changed.

```

class BoundedBuffer from UnitBuffer
  attr prodq buffer
  meth init(N)
    buffer <- {New Channel init}
    prodq <- {New Event init}
    {For 1 N 1 proc {$ _} {@prodq notify} end}
  end
end

```

- Figure 36 A Bounded Buffer Class.

Bibliography

1 Christian Schulte, et. al., 'Oz Standard Modules', the Oz documentation series.

14